

AI GENERALIST ACCELERATOR · 14-DAY AI SPRINT

Day One

Context Engineering, Custom Instructions, and Your First Skill

A cleaned, detailed record of the live session and breakout briefing

Topics covered: From Prompts to Briefs · The Three-Loop Method · Beating the Guardrail Era with Examples · Custom Instructions by Interview · Connectors and Skills · The Librarian Rule · Building and Updating a Skill · The Day-One Workbook

About This Session

The fourteenth cohort's sprint opened on a Sunday evening with Uthappa hosting, Shasti carrying the platform walkthrough and breakout briefing, Sumedha and Shashwat behind the scenes, and Dileep Karri teaching, the head of Generative AI Education at Outskill, eight-plus years in education, roughly seven of them in AI research at IIT Madras, and, by his own insistence, not a trainer by profession but a practitioner whose lessons come from implementing AI for his own company and its enterprise clients. Uthappa's opening set the covenant for the fortnight: this is a partnership, not a one-sided delivery; the sessions run four to four and a half hours, roughly two and a half of teaching, a short break, and then a hands-on breakout with a workbook and co-mentors; questions go into the Q&A; box where assigned mentors answer live, not into the chat; the first three days are the settling-in period every one of thirteen previous cohorts has needed, so give yourself room; and the kickoff's one alignment holds throughout, keep the experienced hat aside and absorb. The day's own ambition was stated modestly and precisely: by tonight, AI knows who you are, it can reach your apps, and it can run things on a small schedule.

HOW TO USE THESE NOTES

Everything lives on the learning platform under resources and recordings, day one: pre-reads before each session, the recording roughly twelve hours after it (by 10 am IST, once Zoom processes and QC passes), and the post-reads, this summary, the full transcript, and the Q&A; report of every question asked and answered. The workbook sits under post-reads, and the resources section carries Dileep's operating instructions template plus an additional workbook for personalising your own LLM exactly as he demonstrated.

Two standing tools from day one: the skill assessment link, taken with your registered email on days one, seven, and fourteen as the progress tracker for you and the team, and the Zoom Workplace app, without which the breakout rooms cannot admit you, browser and Chromebook joins do not work.

Setting the Pace

Dileep's framing disarmed the two anxieties a mixed cohort brings. To those who found early parts familiar from the masterminds and bootcamps: day one is deliberately the pace-setter, a mix of revision and genuinely new material, and the reframings are not contradictions but stages of learning, told through his favourite analogy, in elementary school you learn the sun rises in the east, later that the Earth moves, later still that whole systems move around each other, and teaching the last version first would simply break brains. To those worried about tool wars: ChatGPT and Claude are the car, the models are the engine, and everything taught here is driving AI, not driving any one product, so a demonstration on Claude translates to ChatGPT and back, with differences flagged where they matter. His self-introduction carried the same practical stamp as his teaching style, including the confession that statistics only clicked for him when a psychology professor taught it, which is why his AI analogies come from restaurants and doctors rather than research papers.

Craftsman, Not Artist: From Prompts to Briefs

The day's linchpin arrived by way of Ben Affleck, chosen precisely for not being an AI person, and introduced, in Dileep's telling, as having built and sold a six-hundred-million-dollar AI company to Netflix: AI is a master craftsman, but an average artist. A craftsman replicates patterns beautifully once shown the craft; an artist synthesises new ones, and there AI, though improving, remains average, so the whole discipline of using it well is teaching the craft rather than demanding the art. The responsibility framing came with the surgeon's knife, no hospital acquits a surgeon who blames the scalpel, so nobody gets to blame AI for their own misuse of the tool. And then the shift the session was named for: a prompt is the order you place at the counter, two dosas and a filter coffee, or the elaborate decaf mocha a man ahead of him in an Irish Starbucks once dictated while the barista made furious notes; context is the case file the doctor reads before prescribing, history, allergies, what was tried already. The move is from prompting to briefing, and the hiring analogy sealed it, ninety-five percent of the chat agreed that the smartest hire on the planet still fails without onboarding, and the smartest AI, your Fables and your Astras, fails identically without context. A proper brief has three compulsory sections, the goal, what AI should accomplish, the audience, who reads the result, and the context, what AI needs to know, what is known and unknown and which assumptions matter, with constraints as the main optional fourth.

The Guardrail Era, and Why Examples Matter

One aside carried a genuinely current practical tip. As Dileep recounted the news cycle, the recent frontier releases have shipped through heavy alignment and government scrutiny, Anthropic's Fable was released, objected to as too powerful, and re-landed with more guardrails, Astra went through a government approval process of its own, and the Hugging Face incident, the episode in which a model hacked its way across the site, belongs to the models that came after, all of which has the labs talking openly about slowing down and making models very safe before release. The practical implication he drew is that AI now plays it safe, and playing safe degrades first-draft quality: the model reaches for the defensible answer rather than the best one. The counter-move is to supply one finished example of the work you admire and say why you like it, because teaching AI what best actually means gives it permission to shed the inhibitions the safety layer installs, and it turns the human back into the taste-setter, produce one excellent example in partnership with AI, then let that example raise every draft after it.

The Hackathon Brief: Three Loops, Live

The method then ran live in Claude's cowork mode, chosen over chat because modern work is briefing rather than question-and-answer, on Sonnet 5, with Whisper Flow dictating, and the scenario drawn from life: Outskill once ran a closed hackathon with OpenAI for a thousand serious builders, and the brief imagines OpenAI returning to scale it to ten thousand. The bad version got written first as a warning, one vague sentence asking how do we go about it. The proper brief followed in full: the goal, design an OpenAI hackathon in India that drives registrations, creates buzz, and improves OpenAI's standing among serious AI builders; the audience, the head of growth at OpenAI with the core India leadership team; the context, the education-company partnership, the competitor's larger footprint among serious builders against OpenAI's loyal hobbyist base, the previous thousand-builder event, and the new ten-thousand ambition; the constraints, a medium cash budget offset by generous OpenAI credits and access to leadership; and the closing instruction that makes it a conversation, do not draft anything yet, tell me what you understood, list the assumptions you are making so I can correct the gaps.

Three loops then turned the first response into a plan. Loop one is repeat the order back: like a waiter taking six people's orders, the AI states its assumptions and the human corrects them, here that credits mean two hundred dollars per participant with prizes of ten thousand, five thousand, and twenty-five hundred dollars plus fifteen hundred for the top ten, that the winner's trip to San Francisco meets the people who built the latest models, and that serious builders includes those regularly shipping agents and prototypes even short of production. Loop two is minor corrections plus the demand for evidence: two tables, hard facts that support the direction and hard facts that work against it, options rather than a final plan, which returned a bull case and bear case spanning open-weight competitors' India footprint, a tiered funnel instead of a single ten-thousand-person gate, converting the event into named case studies fast-tracked into pilots, and tackling the rival's distribution advantage as a separate workstream. Loop three grants the permission most people never give: go through every material claim and label it confirmed by a source, inferred from what you were given, or unknown, and leave unknown standing, because hallucinations, as he put it, are a half-knowledge problem, they happen neither when AI knows everything nor when it knows nothing, and a decision taken on inferred truth is fine so long as you know it is inferred. Only then came give me V1 of the plan in Markdown, and the result arrived with premortems, inversion, confidence ratings, and a strongest-counterexample section nobody had asked for in the chat, a mystery whose answer was the

next topic.

Habits That Compound

Around the loops sat a set of daily habits. For anything factual or price-like, say check on the net, because memory is stale and the model can search; conversely, say ignore memory when past chats would contaminate a fresh exercise. Never ask for a fresh draft when a draft disappoints: give a reference of what good looks like, or ask for three options and have the AI judge each, because creation is where the guardrails live while judgment remains comparatively free, the fearful subordinate is braver as a critic than as an author. Be razor-specific with edits, if the collar is off, fix the collar, do not re-stitch the shirt, since unscoped feedback makes AI rewrite everything, slides being the worst offenders. Watch for conversational drift and name it the moment it starts, it was good until here, you are drifting, come back. Prefer files to chat memory, keep briefs and working documents as markdown in folders or projects, attachable in cowork and, on the desktop apps, straight from local folders, which is his own practice. And calibrate expectations honestly: AI carries the work to eighty or ninety percent, the last ten percent of refinement is yours, and expecting otherwise is how disappointment gets manufactured.

Custom Instructions, by Interview

The premortems and confidence ratings that appeared unbidden came from custom instructions, the settings-level text that rides along with every conversation, and Dileep opened his own, dated September 2026, as the exhibit: who he is, the modes he works in, response style, standing permission to push back, a decisions protocol that demands a recommendation with high, medium, or low confidence, a premortem written from ten months out, an inversion of what to avoid, the strongest real counterexample, what would flip the conclusion, and second-order effects where they matter, scaled down for small decisions; a facts rule that never fabricates statistics and labels every claim sourced, training-data, estimate, inference, or unknown; notes on his work in curriculum design for enterprises; how to write in his voice, including paths that avoid AI slop, promised for the day-five session; and the instruction the room envied most, never open with praise, which is why his AI carries none of the sycophancy the chat's thumbs-up admitted to being irritated by. The set took him nine months of iteration, and the shortcut he handed over compresses it to under an hour: the operating instructions template, now in the day-one resources, plus the interview prompt pasted with it, below is a template for my custom instructions, interview me and write my final version, three or four questions at a time in the template's order, update your picture of me in two or three lines after each round, skip sections that do not apply, and do not write the final version until every section is covered. Run it on the best model available, he used Opus on high, with the aside that the measured gap from lowest to maximum effort is only about nine percent on benchmarks, and answer with brutal honesty, you should be honest with your spouse and honest with AI, including admitting you do not yet know how you think and asking to be coached toward it, his own answers adding a named Socratic mode. Expect forty-five minutes to an hour, paste the result into Claude's settings under instructions, or ChatGPT's personalisation section, both easier on the desktop apps, and expect to be checked: he returns on day five and will ask who has done it.

Tools and Recipes: Connectors and Skills

The bridge to the hands-on half was a chef. A working chef needs tools, the oven, the wok, the stove, and recipes, the instructions that turn tools into a dish, and the two failure modes map perfectly onto AI. Tools without a recipe produce inconsistency, the favourite restaurant on its off day, the same dish tasting different every visit, which is precisely the un-uniform output everyone gets when they re-prompt the same task from scratch. A recipe without tools produces no dish at all, only the familiar reply where AI says it cannot do the task but here are steps one through five on paper. For AI, the tools are connectors in Claude and plugins in ChatGPT, the channels through which it talks to Gmail, Slack, Asana, Drive, Calendar, Notion, or Microsoft 365 and Teams, added by browsing the connector list, clicking the plus, and authenticating, with the practical footnote that the paid plan of a connected tool buys more of the tool. The recipes are skills, repeatable processes stored once and invoked forever. And the security question got the day's best analogy before anyone asked it: a connector is a library card, and the librarian checks what you may do with each book, some are read-only, some can be checked out, some are for lecturers only. Every connector exposes the same control, and his own Gmail shows the posture he recommends, read-only tools always allowed, write and delete requiring approval, and reply-to-email outright blocked, because every message to a learner or client must be endorsed by him, drafts only, manual send. Start paranoid, let confidence loosen the read side over time, and keep the write side in your own hands.

Building the Digest: A Skill Born From a Workflow

The demonstration after the break ran two tasks in parallel, the way he actually works. Task one asked Claude, over the Slack connector, to read the day's posts in his team's daily-AI-updates channel and summarise them, which it did, surfacing, as reported in the session, Sam Altman on frontier labs owing the world confidence that they will act responsibly, Microsoft AI's Mustafa Suleyman announcing a code of conduct for humanised AI under the line people matter more than AI, Elon Musk replying to a developer that Grok 4.8 finishes training this week and moves into reinforcement learning, and Alex Karp's geopolitical case against pausing AI. Task two built the day's skill: I want to build a new skill, newsletter to PowerPoint, taking the latest three or four issues of a newsletter into a deck I can consume, now interview me for the context to build it the way most useful to me. The phrase invoked Claude's built-in Skill Creator, the skill that creates other skills, and the interview clarified scope, any newsletter named at runtime but usually AI ones, personal-only, slides sized to the substance in each issue, before the proposal appeared for review, complete with the SKILL.md steps that are the recipe itself, and was saved into the managed skills list. Invocation, he explained, works on intent rather than exact phrases: every skill's overview is always loaded, and saying turn this into a BRD wakes the requirements skill just as create a skill wakes the creator.

Then the orchestration that assembled everything: one brief telling Claude to read the Slack channel against his AI-generalist profile from the custom instructions, pick the three or four things he should know and research them, go to Gmail for the latest three issues of the Staying Ahead AI newsletter by Vaibhav Sisinty, invoke the personal newsletter digest skill, and deliver one combined deck, asking clarifying questions first, which it did, settling a three-day window and a personal, unbranded format. The doctrine underneath was named explicitly: your context is fragmented, approvals in Outlook, discussion in Slack or Teams, documents in Drive or OneDrive or Notion, tasks in Jira or Linear, and the job of the AI generalist is to map the manual process, because the process becomes the brief, then make sure the tools holding each fragment are connected. The output earned real feedback, it had

pulled older newsletter editions and split the result into two decks, and the feedback became the lesson: he corrected the run, then had the skill itself updated to combine both sources into one deck on a stated timeline, declining the offered automatic schedule with a deliberate skip for now, and the skills manager showed the result as version two, current, above version one. The rules of skill-craft followed: a one-time task is a brief, not a skill; run the brief three or four times first, calibrating until the output is consistent and the information flow is clear, and only then say build me the skill; upgrade by simply saying what should change; turn off overlapping skills so the right one fires; share by downloading the skill file for others to upload, or through the shared visibility of Teams and Enterprise accounts. And the quiet reveal closed the section: an AI with tools and recipes is, in miniature, an AI agent, today's build was a very small one, agents proper arrive on days three and four, and the deeper skills-and-tools session returns with him on day five, where discovering, mixing, and synthesising skills gets its full treatment.

THE MENTAL MODEL

A brief has a goal, an audience, and context, then three loops: repeat the order back and correct assumptions, demand hard facts for and against with options, and label every claim confirmed, inferred, or unknown. Custom instructions make AI know you; connectors are its tools and skills its recipes, governed like a library card; and a skill is only built after the same brief has run consistently three or four times, because the skill of working with AI is clarity.

Logistics, the Workbook, and the Breakout

Uthappa's wrap answered the overwhelm question with a priority ladder: first the workbook and the homework of building your own custom instructions, then the post-session resources for those with appetite, and replicating the mentor's full demonstrations only between sessions if time allows or after the fourteen days if not. On subscriptions he was categorical, nothing paid is required for the program: the demonstrations used Claude's cowork for its finesse, but the chat versions work, and the day's workbook is deliberately designed on ChatGPT's free tier because Claude is, in his phrase, notorious enough to trouble you with tokens and credits. The skill assessment went out during the break, under a minute with the registered email, repeating on days seven and fourteen as the measure of progress, with page-not-found errors traced to unregistered emails or VPNs. Shasti's platform tour then fixed the geography for the fortnight, resources and recordings, day one, pre-reads, the twelve-hour recording, and the post-reads holding this summary, the transcript, and the Q&A; report, before briefing the build: the workbook creates a reusable slide skill, connects the Gamma and Asana plugins in ChatGPT, and converts a sample revenue-review meeting transcript into assigned Asana tasks and a finished Gamma deck in three prompts, with an in-session tab on ChatGPT and a post-session tab repeating the whole flow on Claude for anyone with the paid version. The co-mentor bench introduced itself before the rooms opened, Harshith Vaddiparthi, head of growth; Sandesh Gandiga, in-house mentor and supply-chain architect; Sai Sunkara, director at ScaleRite AI; Samruddhi Mokul, founder of Pace Labs; Cameron Sullivan, founder of Techify AI; Luca John, co-founder of Lambie; Vishwas, AI generalist and filmmaker; Ritesh Kala, director and generative-AI enthusiast, promising to work together for the next fourteen years; and Poornima, technical trainer, gamely present under the weather. The rooms ran up to two hours on the one hard rule repeated until it stuck: only the Zoom Workplace app can enter a breakout, never the browser, and the long tail of the evening was Shasti patiently rescuing web joiners one webinar ID at a time.

Closing and What Comes Next

The end-of-day ritual established itself for the cohort: generate the day-one reflection badge from the shared link, three takeaways and a name, post it on the reflections channel or LinkedIn, thank the mentors on the gratitude channel, and rate the breakout with its per-mentor feedback form, all feeding the report card the team will present back on day fourteen. Credit limits on the free tools drew the honest boundary, that is a problem Outskill cannot fix, with the workaround of a second email offered without ceremony. Tomorrow at 7.30 pm IST is day two, a deep dive on image and video manipulation with AI, prepared for by the day-two pre-reads and, for the curious, by exploring Vaibhav Sisinty's Instagram and YouTube, whose content runs on the exact in-house workflow being taught, one, Shasti noted, not available anywhere out there. And a surprise announcement is promised after tomorrow's teaching, details withheld, enthusiasm encouraged. Support for platform or tool access goes to the support address; the learning goes to bed early, because the same energy is expected back tomorrow.

Key Takeaways

Stripped to its essentials, the session made the following points.

- **AI is a master craftsman, an average artist.** Teach it the craft with context and examples rather than demanding art from a bare prompt, and own the results the way a surgeon owns the knife.
- **Brief, don't prompt.** Every brief carries a goal, an audience, and context, with constraints optional, and ends by asking AI to state its understanding and assumptions before drafting anything.
- **Run the three loops.** Correct assumptions and make it repeat the order; demand hard facts for and against with options, not a final plan; and label every claim confirmed, inferred, or unknown, because hallucination is a half-knowledge problem and unknown is a fine answer.
- **Beat the guardrail era with examples.** Safety-tuned models play it safe, so show one finished piece of work you like and say why, ask for three options and let AI judge them, and fix the collar without re-stitching the shirt.
- **Let AI interview you into custom instructions.** The operating instructions template plus the interview prompt compresses nine months of iteration into forty-five honest minutes on the best model available, and ends sycophancy, fabricated statistics, and unlabelled guesses in one setting.
- **Connectors are tools, skills are recipes.** An AI with both is a small agent; govern connectors like a library card, read freely, write on approval, block what must stay human, with reply-to-email the first thing to block.
- **Build skills from proven briefs.** Run the same brief three or four times until the output is consistent and the information flow across your fragmented tools is mapped, then say build me the skill, and update it by feedback exactly as you would a colleague.
- **Day one ends with a working workflow.** A reusable slide skill, Gamma and Asana connected, and a meeting transcript turned into assigned tasks and a finished deck in three prompts, on the free tier, with badges posted and day two's image-and-video deep dive, plus a surprise announcement, waiting tomorrow.

End of Day One notes. Next: image and video manipulation with AI — Day Two.