

## AI GENERALIST ACCELERATOR • 14-DAY AI SPRINT

Day Two

# Image and Video with AI: Clones, Footage, and the Content Playbook

*A cleaned, detailed record of the live session and breakout briefing*

Topics covered: Why Content, and the Four Blockers • The Five-Step Playbook • Finding Ideas • Hooks and the 80-20 Rule • Writing DNA and a Script Skill • Building Your Clone • The Master Director Formula • Editing and Assembly • Bonus Office Hours • The Aura Beans Workbook

## About This Session

Day two handed the evening to Uttam Gupta, founder of the AI transformation company AgentValley.ai, formerly director of growth at Outskill, and the person who built Outskill and GrowthSchool with Vaibhav from day one, raising five million dollars, scaling the team from zero to two hundred and fifty, and helping grow the world's biggest AI community of over ten million people that this cohort now belongs to. He teaches what he runs: a hundred and seventy thousand followers of his own across LinkedIn and Instagram, a hundred thousand people trained across corporates like Reliance, Upgrad, and Viacom, and, most relevantly, the in-house playbook behind Vaibhav's channels, where, as Shasti framed it, ninety-nine percent of what you see is AI-generated, a workflow she was clear is not available anywhere out there. Her three ground rules held for the night: the taught half uses paid tools while the breakout uses free ones, so watch and learn now and execute against the recording later; the process normally takes the internal team a full day and was compressed into three hours, so lean on the recording, the Q&A; box, and the AI companion rather than panic; and an announcement was promised for after the teaching, delivered on schedule and covered below.

### HOW TO USE THESE NOTES

Everything the session used is on the learning platform under day two: the cheat sheet holding every hook formula, writing framework, and the Master Director prompt menus, Uttam's full deck, and, crucially, the additional workbook titled Instagram AI Workflow, which converts his entire paid-tool playbook into step-by-step form for anyone who wants to run it post-session. The in-session workbook is the free-tool Aura Beans commercial described at the end.

Breakout rooms are never recorded, a Zoom limitation, so a missed workbook is solved solo with help from the Ask Doubt channel.

## Why Content: The Four Blockers and the Multiplier

---

The opening was a test: two reels, a family-salon-to-thirteen-hundred-crore perfume-brand story and a thought-reading wearable pitch, with the chat voting on which was AI. The room split every way, and the answer was both, entirely AI-generated, millions of views between them, which set up the founding observation of the whole playbook: audiences care about the value of what you say, not whether a camera was present. Four questions then found the four blockers, no time, can't write, no ideas, and fear of judgment, each drawing a wall of me's from the chat, and the session's promise was to dismantle all four with systems. The argument for bothering at all was the one every professional recognises: people with less experience who earn more opportunities do so not because they know more, but because more people know what they know, and being good at your work used to be enough, but now you must be known for what you know. What follows from publishing, he insisted from experience, is opportunities you never applied for, clients who arrive pre-sold on your credibility, a reputation that introduces you before you speak, and leverage, content as real estate, an asset compounding while you sleep, none of which requires becoming a content creator in the cringe sense, only learning to talk about what you already know. The proof stack was personal and transferable: over six hundred million views and thousands of leads generated for him and his clients, Vaibhav grown from a hundred thousand to over three million followers, the process sold as a service at three thousand dollars a month and up, and, for anyone starting from nothing, Definay, an Outskill learner from Turkey with eight hundred followers and no time, reaching fifty-three thousand followers in six months without appearing on camera, a TEDx stage, an AI-training career, and over two hundred thousand dollars in income.

## The Five Steps, and the Split

---

The playbook has five steps: find ideas worth creating, script them into a format people actually listen to, produce the video without a camera by building a clone of yourself, make it engaging with AI footage, and assemble everything in the edit. The honesty was in the percentages. Ideation and scripting are eighty percent AI with twenty percent human judgment; production is a hundred percent AI, he has not shot a single video in twelve months while his account kept growing; footage is a hundred percent AI; and editing inverts to only twenty percent AI with eighty percent human craft. Deliberately not fully automated, because the human twenty percent, the decisions, the perspective, the taste, is what keeps a Vaibhav in Vaibhav's content even with a team and a process running underneath, and it is precisely what makes the output yours rather than everybody's.

## Step One: Six Places Ideas Come From

---

Ideas come from six sources, all demonstrated. Perplexity, run as his browser, for research and staying current. Twitter, which he was blunt about, whatever reaches LinkedIn, YouTube, or Instagram appears on Twitter first, the decision-makers live there, so consume it, doom-scroll it, and mine its hashtags. Reddit, for the private communities of any niche, real estate to healthcare, where real questions and pain points surface. Claude, his Mr. Claude, for deep research, news, and later writing, with the practical aside that a daily-news prompt needs no automation platform at all, the built-in scheduled tasks run it every morning at nine and can email the result. Competitors, four or five listed and followed, studied for what works rather than copied. And trends, the things repeating on your feed and in the news, the fastest route to reach because the demand already exists. The live run asked Claude for the last forty-eight hours of AI news, and the week's biggest story arrived, told in the session as follows: Anthropic's Dario published an essay saying AI is moving too fast and too dangerously and asked the industry to slow down; Sam Altman agreed and retweeted, Elon Musk agreed that Dario is right, rivals suddenly on the same side; then Donald Trump refused, declaring the US will not slow down because China would get ahead, AI being bigger than weapons; China called the whole exchange fearmongering, a trap to keep America ahead, and chip stocks crashed on Monday. Picking from a list of angles comes down to three filters, impact, timing, and whether you can genuinely add perspective, and one further prompt collapsed the chosen story into a two-line topic ready for scripting.

## Step Two: Hooks First, Everything Else Second

---

A reel has three parts, hook, body, and call to action, and a sixty-second reel is roughly a hundred and fifty words, the sweet spot for attention spans outside YouTube. The rule that surprises everyone is the time split: eighty percent of writing effort goes into the hook, the first five seconds, and only twenty into body and CTA, for three reasons hammered in turn, viewers decide within the first sentence, the algorithm amplifies whatever holds attention, and the best body in the world is worthless if nobody stayed to hear it. The frameworks then arrived in layers, all preserved in the cheat sheet. The three-step hook formula: establish the topic, deploy a contrast word, but, however, here's the interesting part, as the stun gun, then snap back with a sentence that runs opposite to the opening. Six hook archetypes, each worked against the day's story and the room's industries: the fortune teller, this changes AI forever; the experimenter, I tried X for thirty days, here's what happened; the teacher, here's how X solved Y, let me explain; the magician, a striking visual or sound that forces attention; the investigator, I found something most people don't know; and the contrarian, you're doing X wrong, here's why. On top of those, the extra styles, pattern interrupt, delete your ChatGPT, here's why, bold claim, this one tool replaces your entire video team, curiosity gap, Google just did something that changes everything, direct challenge, you're doing X wrong, here's the proof, and future shock, in six months every brand will use this and most aren't ready. The skill, he stressed, outlives content: the same psychology opens his trainings, structures his presentations, and is exactly what an email subject line is, and every hook should pass the checklist, does it create urgency or curiosity inside five seconds, speak directly to the viewer, imply a transformation or reveal, and carry a visual hook where possible.

## Borrowed Voices, and the Writing DNA Skill

---

For the body there are two routes, write the way you talk, or, for beginners without a reel voice of their own yet, borrow one, and he spent real time dissolving the guilt around that: you learned walking, writing, and driving by copying, every writer, comedian, and musician started in imitation, Steal Like an Artist wrote the book on it, and the borrowed voice is scaffolding you grow out of, not theft. The method is five steps run live on Vaibhav. Pick one or two creators whose expression fits you, they need not be the biggest, only great at expressing. Find their top ten reels with ViewRank, a free Chrome extension that, one click on any Instagram or TikTok profile, scrolls and scrapes fifty to five hundred reels and ranks them by views, likes, or comments, with a copy-for-Sheets button that lands the whole dataset in a spreadsheet, the harder Apify route acknowledged and set aside. Transcribe the winners with HappyScribe, free for about ten minutes of video by pasted link, or weighin.ai, which gives a few hundred free minutes and swallows long YouTube videos whole. Collect the transcripts into one inspiration document, the raw material. Then let Claude do the learning: the document went in with a request for a full writing analysis, a writing DNA, and out came the anatomy of virality, the hook moves, the important price number, the villain move, show-and-tell, the disbelief opener, the structure of shock hook into doubt-killer into quick context, the retention devices, the pacing, and the vocabulary itself, crazy, insane, wild, genius, mind-blowing, master stroke. The analysis became a mega prompt in the CoStar framework, and the mega prompt became a skill, the Vaibhav viral writing DNA skill, built through Claude's skill creator complete with reference documents and an audit checklist, do it once per creator and the recipe is permanent. Script generation then took seconds: the news research pasted in, the skill invoked, a first draft out, and the step most people skip made explicit, a hundred percent AI script sounds like nobody because it sounds like everybody, so feedback loops trimmed it to a hundred and fifty words with five hook variations and more spice, and his own perspective was dictated in, that AI cannot be stopped once the weapon is out, so the answer is conscious users, policies, and protective institutions rather than pauses, producing a final script that carries him, with the confession that in real production he still edits scripts by hand.

## Step Three: Building the Clone

---

The clone has two halves, and the priorities are unequal: audiences forgive an imperfect visual but never a robotic voice, so audio must sit at ninety-five percent plus while video can survive at ninety. Two paths were laid out. The beginner path lives entirely in HeyGen, worth its twenty-nine dollars a month for unlimited video generation against the free tier's three: create a digital twin by uploading a video, fifteen seconds is the stated minimum but one to two minutes is the real recommendation, complete the on-camera consent step, and train; polish the voice by uploading, again, far more than the minimum, ten to twenty minutes of audio; and, the feature that removed a whole category of excuses, add a look, where design-with-AI generates new outfits and settings, a full podcast-studio version of him appeared without a camera ever being lifted, one video-trained avatar per paid account but unlimited looks, with image-trained avatars serving his clients. The recording rules that decide everything came next: phone camera is fine, natural window light or a cheap clip-on, hand gestures generic and below the face so lip sync survives, expressions neutral rather than committed to one emotion, and never a moving background. For voice, the mistake half the room admitted to planning, reading a book aloud, was banned along with podcast audio, the exact error he and Vaibhav made two and a half years ago that produced slow, boring clones: record reel-style scripts, AI-generated if needed, with the energy, pace, and expression you want back out, garbage in, garbage out, in parts if necessary. The advanced path swaps audio to ElevenLabs, toured in full: voice design conjures any synthetic voice from a prompt, the instant clone needs ten to thirty seconds and was demoed live to prove its limits, recognisably him but with a borrowed accent, fine for narration, never for a personal brand, and the professional clone demands thirty minutes of clean audio and occupies the single professional slot each paid account allows, one account, one voice. Then the honest economics: a month ago he moved his own volume, thirty to fifty videos monthly, from ElevenLabs, state of the art but fifty to eighty dollars a month at his usage, to Fish Audio, the same workflow at roughly a fifth of the price, fifteen dollars for a quarter-million credits, while Vaibhav stays on ElevenLabs, and the professional-clone playback of the Trump script was indistinguishable from the man presenting it. Production closes the loop in HeyGen's scene editor, script pasted or audio uploaded, avatar and voice switchable on the right, and a model choice explained with the day's best analogy, Avatar 3, 4, and 5 are HeyGen's Opus-and-Fable ladder, the newest is best and most expensive, the paid plan giving six hundred credits of Avatar 5 against unlimited Avatar 3, so 5 is reserved for the videos where gesture accuracy and emotional adaptation matter and 3 carries the volume; lip sync renders only on generate, five or six minutes later the reel existed, played back to the room with a Zoom-lag disclaimer and a chat full of boom.

## Step Four: AI Footage and the Master Director Formula

---

The gap between a raw clone video and the ones on his and Vaibhav's feeds is B-roll, the supporting clips and images that keep a video from being a talking head, and, usefully, that absorb attention so completely that clone imperfections vanish, audiences do not scrutinise the way students of the process do, and even Vaibhav's real footage now draws is-that-your-AI-clone comments. The tool for all of it is Higgsfield, a marketplace that bundles the frontier image and video models, Nano Banana Pro, GPT Image, Kling, Seedance, Veo 3, under one subscription instead of ten, a company he noted has run from twenty million to seven hundred million dollars of revenue in a year. His thumb rules cut the choice paralysis: images on Nano Banana Pro or GPT Image, image-to-video on Kling 3.0 or Seedance 2.5, dialogue video on Veo 3 or Kling 3.0. Then the framework that makes non-creatives dangerous, the Master Director prompt formula, seven elements, subject, action, environment, camera, lens, lighting, and style reference, with the cheat sheet supplying menus for every jargon layer, camera angles from low to bird's-eye, movements like dolly-in and rack focus, lenses from 24 to 85 millimetres, lighting from golden hour to edge light, film stocks like Kodak Portra 400. The proof was a before-and-after: women in a cafe produced four generic variations for fourteen credits, while the full formula, a young Indian woman in a cream linen blazer sipping from a ceramic espresso cup at a marble table in golden-hour window light, shot on 85 millimetres with shallow depth of field, slight Dutch tilt, Portra 400 film grain, produced exactly the image in his head, same tool, same model, completely different output, because the words are not decoration, they are how imagination survives translation. Image-to-video followed, the cafe image animated in seven seconds of gentle zoom for fourteen credits, with the cost law stated plainly, more seconds and higher quality cost more credits. Veo 3 then added dialogue, a generated character delivering I think I might lose my job, not today, maybe not tomorrow, but because of AI it feels possible, and, more startlingly, the same clip regenerated from a single photo of Uttam himself, his face, near enough his voice, no studio. For the day's actual reel he showed the working method: Claude was asked what hook visuals would stop the scroll, returned image prompts, and Higgsfield turned them into an angry pointing Trump, the three rival leaders holding hands, a combined scene of Trump shouting at them with the US and China crumbling behind, plus the asset list for the body, giant chess kings on dark marble, a futuristic speedometer, an emergency alarm, a China-versus-US frame, and a warm studio background generated for the edit. The bonus toy was motion control, upload any dancing or talking video, supply a generated photo, a cute AI child, a Pixar version, yourself, and Kling transplants the performance onto the new subject, the mechanic behind every meme reel of dancing kids and animals, and equally the way to replace yourself in any recorded clip.

## Step Five: Assembly

Editing is where the split flips to eighty percent human, and the tour ran in VEED, with CapCut, SubMagic, and Edits named as equivalents. Five elements make the edit: the hook clip up front, music, B-rolls, subtitles, and background. The walkthrough did each in minutes, resize to vertical, remove the clone's background, add a music bed and duck it under the voice, auto-generate realistic B-rolls from the tool itself and layer the Higgsfield assets on top, shrink the talking head and drop the generated studio backdrop behind it, then the highest-leverage click of the section, auto-subtitles with the viral style presets, the punchy animated captions every winning reel wears. His disclaimer was honest, five minutes of editing where thirty would polish, but the final render played end to end, hook to Trump's line that the only guardrail AI needs is a strong president, to China's fearmongering charge, the Monday chip crash, and the closing turn, the people building AI are scared, the people controlling it are not, and your job is stuck right in the middle, comment below, follow before the next move drops. Idea to published-grade reel in roughly two hours forty against the team's usual twenty-four, and the challenge attached: before the week ends, make something, a clone, a clip, the fun version, and post it or send it to him, because with time, writing, ideas, and judgment all systematised, the only step left is publishing.

### THE MENTAL MODEL

Five steps, one split: ideas and scripts are eighty percent AI and twenty percent you, production and footage are fully AI, editing is mostly you. Spend eighty percent of writing time on the first five seconds, train clones on energetic reel-style data rather than books and podcasts, prompt images like a director with all seven elements, and let B-roll carry the attention your clone doesn't have to. The process is simple, not easy, and the human twenty percent is what makes it yours.

## The Announcement: Bonus Office Hours

The promised surprise turned out to be more teaching. Moved, Shasti said, by the cohort's first-two-days commitment, including one learner who compiled the chat and Q&A; into an FAQ document and posted it to the reflections channel for everyone, the team is adding Bonus Office Hours: a dedicated one-hour session every weekday from tomorrow through Friday, day three to day five, from 6 to 7 pm IST, where a mentor answers questions live, not just on day one and two content or a stuck workbook, but anything about AI at all. The mechanics are deliberately simple, everyone stays muted, questions go into Slido exactly as introduced in the breakout rooms, the joining link appears as its own live session on the platform's live-sessions tab, WhatsApp reminders go out, and every hour is recorded to the platform for anyone who cannot attend. It is explicitly optional, with none on the weekend's full-day sessions, and the daily rhythm from tomorrow is office hours 6 to 7, a break, then the main session from 7.30 running to 11.30 or midnight depending on the breakout.

## The Aura Beans Breakout, and Logistics

---

The breakout deliberately did not replicate the paid-tool playbook, which needs a day and subscriptions; instead the room built its first AI commercial for Aura Beans, a coffee brand invented for the exercise, entirely on free tools: assets generated with Nano Banana, animation in Google Flow, and assembly in Canva, the thirty-minute AI commercial challenge meaning a five-to-ten-second ad produced in about half an hour. The known constraint was Google Flow's fifty free daily credits, met with the standing advice, stick to the workbook's given prompts, aim for learning rather than perfection, and if credits run dry, a second email account is the workaround Outskill cannot provide but can recommend. For everyone hungry for the real thing, the additional workbook on the platform's day-two resources, Instagram AI Workflow, is the entire three-hour playbook converted to steps for post-session use, alongside the cheat sheet and the deck. Jay, founder of Jay Prince, and Jordan, agency co-founder, joined the returning co-mentor bench for the rooms, which ran the usual app-only rules. The wrap set tomorrow's double rhythm, office hours debut at 6 pm IST, then the main session at 7.30 opening the world of automations, with the homework to explore n8n ahead of time and read the day-three pre-reads, and finished videos posted to the reflections channel where the team will actually see them.

## Key Takeaways

---

Stripped to its essentials, the session made the following points.

- **Be known for what you know.** Both opening reels were AI, audiences care about value rather than cameras, and the four blockers, time, writing, ideas, and judgment, are systems problems this playbook solves, with opportunities, trust, reputation, and leverage as the compounding returns.
- **Five steps, one honest split.** Ideas and scripts run eighty-twenty AI to human, production and footage are fully AI, twelve months without a camera, and editing stays eighty percent human, because the twenty percent of judgment and perspective is what keeps the content yours.
- **Ideas are a routine, not a muse.** Perplexity, Twitter first, Reddit's niche communities, Claude with a scheduled daily news prompt, four or five tracked competitors, and trends, filtered by impact, timing, and whether you can add perspective.
- **Spend eighty percent of writing time on five seconds.** Hook, body, CTA in a hundred and fifty words, the three-step formula and six archetypes to draft hooks fast, and the same psychology powering presentations and email subject lines forever.
- **Borrow a voice, extract its DNA, make it a skill.** ViewRank scrapes the top reels, HappyScribe transcribes them, Claude distils the writing DNA into a mega prompt and a permanent skill, and the hundred-percent-AI draft that sounds like nobody becomes yours through feedback and your stated perspective.
- **Audio is the clone's soul.** Viewers forgive imperfect video, never a robotic voice, so train on ten-plus minutes of energetic reel-style speech, never books or podcasts, use HeyGen for the twin and looks and a one-per-account professional clone for the voice, routing volume through cheaper audio once quality is proven.

- **Prompt images like a director.** Seven elements, subject, action, environment, camera, lens, lighting, style, turn generic outputs into your exact imagination on Higgsfield's one-subscription marketplace, and B-roll absorbs the attention your clone's flaws would otherwise draw.
- **Publish this week.** The edit is five elements in VEED with viral subtitles, the full reel shipped in under three hours against the team's usual day, office hours start tomorrow at 6 pm IST for every question, and the only remaining step the playbook cannot do for you is posting.

*End of Day Two notes. Next: automations — Day Three.*