

AI GENERALIST ACCELERATOR · 14-DAY AI SPRINT

Bonus Office Hours One

Office Hours One: Portable Memory, Tool Discovery, and Learning Anything

A cleaned, detailed record of the first Slido question hour

Topics covered: Monetising AI Content · The Cloning Stack, Clarified · Your Memory on GitHub · Picking and Finding Tools · Tokens and Local Models · Marketing Workflows and Scraping · NotebookLM · Skills, and Learning Anything in Six Months

About This Session

The bonus office hours announced on day two opened on schedule at 6 pm IST, an hour before day three's main session, with Sumedha hosting and Harshith Vaddiparthi, Outskill's head of growth and a breakout-room regular, in the answering chair. The format is now the standing one: everyone stays muted, questions go into Slido through the link in chat, the thumbs-up upvote decides the order because popularity ranks the queue, not the mentor, and anything unanswered carries automatically into the next day's hour, which is exactly where the session ended, thirty-six questions still standing when time ran out. Harshith widened the brief in his first minute, the hour is not restricted to session content, anything about AI is fair game, and drew the one boundary that recurred all evening: dashboard, LMS, badge, and course-feedback questions belong with support at outskill.com or the tech-support channel, where the team can actually act on them.

HOW TO USE THESE NOTES

Every link Harshith shared landed in the chat and travels with the recording: Product Hunt, There's An AI For That, the Token Optimizer repository, Ollama, NotebookLM, Apify, Postiz, Resend, OpenRouter, and skills.sh. The recording of every office hour goes to the platform like any session.

The single most reusable tip of the hour doubles as the study method for the whole course: build a NotebookLM notebook called Outskill Course, upload every worksheet, resource, and note as it arrives, and question, summarise, or podcast your own materials without spending a credit.

On Content, Monetisation, and the Cloning Stack

The first cluster of questions trailed day two's video session. Does monetisation favour human shoots over AI content? It favours neither, platforms pay for value, not provenance, ninety percent of what floods every feed already carries AI somewhere in its pipeline, and Harshith's own newsletter is the working example, AI scrapes the news and drafts the issue, he edits and adds his pointers, and sponsorships pay him all the same. Why does the playbook use ElevenLabs when HeyGen clones voices too? Because tools have centres of gravity: HeyGen is the best video-cloning and lip-sync tool and its voice feature is merely good, ElevenLabs is the best pure audio platform, Fish Audio surged out of nowhere on price, and Harshith's personal stack goes one step further, an open-source Qwen Audio 3 model running on his own GPU, with the practical note that a voice cloned in ElevenLabs imports into HeyGen anyway, so learn the extra tool for the use cases that never need video. The picture-training question got its proper name, a LoRA: gather roughly twenty photos of a client or character, train the model on that folder, and every output thereafter features that person, exactly the technique behind the client pages people had noticed. The Google Flow aspect-ratio confusion from the breakout was untangled, the images were already generated in Gemini so the image setting is irrelevant, only the video-generation default, set to 9 by 16, matters, garnished with his own blooper reel, yesterday's generation gave a character two right hands. And the workbook-format worry, twelve Gamma blocks instead of a deck, drew the general law: the output follows the ask, request a PowerPoint and you get one, tag Gamma explicitly with the at-shortcut when you want Gamma, and say twelve slides when you mean twelve slides. On the recurring paid-tools feedback he was candid twice over, sessions do flag paid tools upfront, and the working economics of his own client business are worth copying, the client pays for the credits and the tools, you are paid for the skill and the taste, so master the skill on free tiers and let tools be somebody else's line item.

On Portable Memory, and GitHub for Non-Coders

The best question of the hour came from a learner who has fed ChatGPT more than a year of context and fears the lock-in: is there shared memory, so a move to Claude does not start from zero? Harshith's answer was a live demonstration of the principle he framed as your information stays on GitHub, and your tools can be anywhere. He asked his ChatGPT to summarise what it knows about him professionally, then, with the GitHub plugin tagged, told it to create a folder on his GitHub called Knowledge Base and push the summary there, done in one prompt. The point is ownership: GitHub is yours, so tomorrow's tool, Claude, DeepSeek, whatever ships next, simply connects the same plugin and reads the same knowledge base, no vendor holding your memory hostage. Around it he dismantled the GitHub-is-for-coders assumption from personal experience as a self-declared non-coder: repositories are just storage for code and documents, free to create, uploadable through the plugin from ChatGPT or Claude with the single verb push, and his own website is the proof of the whole loop, harshith.com vibe-coded entirely with Fable, its code living in a private repository he could make public or even sell if he chose. The companion question about all-in-one subscription bundles got the honest non-answer it deserved, it depends on your workflow, test tools individually first and buy bundles only when the package genuinely maps to what you use.

On Picking Tools, and Finding New Ones

The how-do-you-choose question, ChatGPT or Claude or Gemini or Meta's just-launched Muse, received the answer the course has been building toward: stop hyperfixating on tools and own the methodologies. ChatGPT's best model today is GPT-6, Claude's is Fable 5.1, and neither fact matters much, because meta-prompting, MCP servers, and plugins are skills that survive every release, GPT-7 included, and the practical procedure is always the same, look at your existing workflow, find where the manual hours go, and ask ChatGPT to help you build a clear-cut automation for exactly that. For the perennial give-us-a-top-five-tools request he refused the list and handed over the two discovery engines instead: Product Hunt, where everyone from indie hackers to billion-dollar companies publishes launches, Google's Gemini 3.8 had appeared there that very day, and There's An AI For That, the searchable directory that proves its own name, pop-ups and all. The multi-model platforms question, is something like GetMulti.ai legal and economical, got both answers, perfectly legal, genuinely efficient for comparing models side by side, with OpenRouter named as the heavyweight alternative, one API that switches between any model, fresh, as he told it, off a billion-dollar acquisition by Stripe two years into its life, and DeepSeek recommended without embarrassment as the cheap-and-smart option in any router.

On Tokens, Local Models, and Tonight's n8n

Token anxiety produced the hour's most concrete stack. For optimising usage rather than endlessly rewriting prompts, he shared the Token Optimizer repository with an installation ritual that carried its own lesson, do not run it in the web chat, install the ChatGPT desktop application first, echoing Dileep's day-one advice that the desktop apps see more of your computer, then paste the repository link and say install this. For watching consumption, Claude's own settings expose sessions and tokens burned, and the Mac menu bar has a family of free open-source trackers, Context Bar and its codexbar cousins, showing daily and weekly limits at a glance, with his cheerful admission that he had deleted his and should reinstall it. The local-LLM question got the full recipe: go to Ollama, browse the models, start with something like Gemma, and size against your machine, his own MacBook's eighteen gigabytes of memory runs anything under roughly twelve billion parameters, and if the terminal command intimidates, paste it into ChatGPT with your configuration and ask how do I install this on my Mac, because, in the refrain of the hour, just use ChatGPT for everything. The pre-session question, what setup does tonight's n8n class need, took ten seconds, an account at n8n.io, whose fourteen-day free trial covers the entire course. And the enterprise-security question, can n8n be sold into banks and government, drew a confident yes built on the same open-source fact the evening session would teach an hour later: a company valued, as he put it, around fourteen billion dollars, its whole codebase public on GitHub with two hundred thousand stars, self-hostable on your own servers so nothing depends on the cloud version, credentials he backed with his own US enterprise clients, a just-closed signage company running Microsoft Enterprise, and the n8n instance he hosts on his own website.

On Marketing Workflows, Scraping, and Learning Anything

The marketing-stack question opened his own toolshed: Postiz schedules content across some thirty social platforms from one place, Resend sends his newsletter programmatically to the subscriber list, and the same GitHub integration that stores his knowledge pushes each issue onto his website as a blog post, the theme being combination, wire the tools together rather than admiring them separately. Scraping Instagram and TikTok for niche content got the goldmine answer, Apify, described as an app store for APIs, where the Instagram scraper alone serves four hundred thousand users, paste a profile, collect the data, and, once n8n clicks, wire Apify into workflows, or simply hand the link to Claude or ChatGPT. The how-do-I-study-this-course question produced the hour's longest demonstration, NotebookLM: create a notebook, name it Outskill Course, upload every worksheet, resource, and lecture note as a source, and the notebook becomes a private companion that answers questions only from your materials, builds slide decks, renders audio and video overviews, and, the party trick he played live, hosts an interactive podcast you can interrupt, he joined the generated conversation mid-sentence to ask about a report's first-quarter numbers and the hosts answered him by name, all free of credits as a Google product. The deepest answer went to the learner asking how AI can teach unfamiliar enterprise stacks, PySpark, SQL, AWS, Snowflake, Unix: he quoted the week's viral line from a thirty-year Salesforce enterprise veteran, the difference between a newbie using AI and the most proficient person in enterprise AI is just six months, told his own version, closing an enterprise client whose Coldbridge system he had never touched and serving them through ChatGPT plus a browser-automation agent, and then showed the mechanism the cohort will formally meet soon, skills.sh, a marketplace where practitioners package their expertise as installable folders of system prompts: search Snowflake, find a semantic-views skill, copy its command into Claude or ChatGPT, confess honestly that you have never heard of the tool, and let the installed skill plus the model make you dangerous in weeks, the same route he is taking into 3D printing, a Bamboo Lab A1 on order, GPT-6, and a Blender skill standing in for the modelling experience he does not have. Two quick definitions closed the hour, an API is a waiter, carrying your order from the front of house to the kitchen and returning with the result, an API key the unique pass that lets you order at all, created in minutes at the ChatGPT API platform; and his parting curiosity, Google has open-sourced a fly's brain map, which he has handed to his own agent to build on, the intersection, as he put it, of technology and biology, with a blog post to prove it.

THE MENTAL MODEL

Own your memory on GitHub and let the tools be interchangeable; own the methodologies, meta-prompting, MCPs, plugins, and let the models race each other; discover through Product Hunt and There's An AI For That rather than waiting for lists; study your own materials in NotebookLM; and when a skill is missing, install someone else's from skills.sh and close the six-month gap.

Key Takeaways

Stripped to its essentials, the hour made the following points.

- **Value monetises, not provenance.** Platforms pay for what content delivers, not whether a human shot it, and the working client economics put credits and tools on the client's bill while you charge for skill and taste.
- **Keep your memory where you own it.** One prompt dumps ChatGPT's knowledge of you into a GitHub Knowledge Base folder via the plugin, and from then on any tool, Claude or whatever ships next, reads the same base, no lock-in, no coding required.
- **Methodologies outlive models.** GPT-6 and Fable 5.1 are today's bests and tomorrow's history; meta-prompting, MCPs, and plugins are the skills that transfer, and every automation starts from where your own workflow bleeds manual hours.
- **Discovery is a habit, not a list.** Product Hunt and There's An AI For That surface the daily launches, GetMulti-style platforms and Stripe-acquired OpenRouter compare and switch models on one API, and DeepSeek stays the cheap-and-smart pick.
- **Watch and shrink your tokens.** The Token Optimizer repository installs through the desktop app, Claude's settings and free Mac trackers like Context Bar show the burn, and Ollama runs Gemma-class local models sized to your RAM, under twelve billion parameters on an eighteen-gigabyte machine.
- **NotebookLM is the course companion.** Upload every worksheet and note into one notebook and question, summarise, deck-ify, or podcast your own materials, interrupting the generated hosts live, all without spending a credit.
- **Any stack is six months away.** The Salesforce veteran's line holds: with skills.sh packaging expert system prompts for Snowflake, Blender, or anything else, honest prompting, and an agent, a newbie closes on thirty years of experience, and the API standing between you and any tool is just a waiter with a key.

End of Office Hours One notes. Next: Office Hours Two at 6 pm IST.