

AI GENERALIST ACCELERATOR · 14-DAY AI SPRINT

Bonus Office Hours Two

Office Hours Two: Connecting Everything, Running Locally, and Building Live

A cleaned, detailed record of the second Slido question hour

Topics covered: MCPs for Non-Coders · Confidential Data, Locally · Enterprise Integrations · Copilot and Azure · Connecting GitHub · Grok and Mixture of Experts · NotebookLM Limits · A Council of Experts Built Live · Hyperframes Video · The skills.sh Hack

About This Session

The second office hour ran the standing format, Sumedha hosting, Harshith Vaddiparthi answering, Slido ranking the queue by upvotes, this time opening on the questions carried over from day one before reopening the queue ten minutes in for fresh ones. Harshith arrived with a new haircut and, he confessed, a real face, yesterday's avatar having been an AI version of himself, which set the tone for an hour that was less lecture than live workshop: he connected MCPs, installed local models, wired GitHub, and, while answering other questions, had an agent vibe-code, host, and publish a working multi-model comparison app in under five minutes. The one boundary held firm throughout: anything about course structure, training hours, or the agenda goes to support at outskill.com, because a mentor can answer AI questions but not programme ones. He closed with a promise to have his agent compile everything demonstrated into a document for the platform, and with the remaining questions rolling into tomorrow's hour.

HOW TO USE THESE NOTES

Links shared in chat and travelling with the recording: the Linear and Notion MCP documentation, the n8n GitHub repository, Ollama and Gemma 4, GetMulti.ai, OpenRouter, Hyperframes, skills.sh, and Harshith's Azure-provider diagram screenshot. Every office hour is recorded to the platform like any session.

The recurring move of the hour is worth stating once: whatever the tool, find its documentation or MCP link, paste it into ChatGPT or Claude, preferably the desktop app, and say connect to this, install this, or look into my configuration and recommend what to do. You do not need to know how; you need to know where the link is.

On MCPs for Non-Coders, and Working Locally

The opening question came from a learner whose company had offered a shift into an AI role, automation, CRM, project management, employee management, with no coding background, and Harshith's answer was an unqualified yes with one skill to master, plugins, which underneath are MCP servers. The demonstration made it mechanical: for a project management tool like Linear, search Linear MCP, copy the link, paste it into ChatGPT or Claude, and say connect to Linear MCP; the model reads the documentation, connects, and hands you an authentication link for your company account. Notion works identically through its official MCP documentation, and once connected the instruction is just conversation, look into my Gmail, check the notes from my project manager, and update my Notion. The confidential-data question got the local stack in two halves. For the workflow tool, n8n's GitHub codebase is public, so copy the repository link into ChatGPT and say install n8n locally and it downloads and sets it up on your own machine. For the model, Ollama is the app store for local LLMs; he recommended starting with Google's Gemma 4, pasting the link into the ChatGPT desktop app and asking it to check his Mac's configuration and recommend what to install, which it did, reading his M3 Pro with eighteen gigabytes of unified memory and a fourteen-core GPU and proposing a version that would run comfortably. The security question that followed later got the same answer drawn as a diagram: a downloaded model like Gemma or Qwen does its entire computation inside your own chip, nothing leaves the machine, and the model is not learning from your data because it is a static copy you own, which makes local LLMs the most secure option that exists.

On Enterprise Integrations, Copilot, and Azure

Asked for a real example of AI integrated into enterprise systems, ERP, industry applications, infrastructure, rather than sales and marketing, Harshith described his current client engagement on Salesforce and Coldbridge, a platform he had never heard of before the project, and his method, which is the method for every unfamiliar tool: ask ChatGPT to do a complete run-through of the platform, its features, pricing, partners, documentation, and build yourself a resource, then find every integration and API the system exposes and hand those to the model to build on, with QuickBooks, Outlook, and NetSuite named as enterprise platforms that all publish integrations, and a note that Salesforce has a Claude integration in beta. The learner whose organisation permits only Microsoft Copilot got both a redirect on the course-content question to support and a genuinely useful insight drawn as a diagram: Copilot runs the GPT models, because Microsoft's Azure is OpenAI's cloud provider, so the same models sit underneath ChatGPT and Copilot, every prompting methodology from the course transfers unchanged, and Copilot's only real difference is that the Microsoft 365 integration comes native where ChatGPT and Claude need it added. A related question, automating Outlook 365 replies into drafts, got two routes: in n8n, an Outlook node granted access to create drafts and send email, or, far simpler, Claude's single Microsoft 365 connector that exposes SharePoint, OneDrive, Outlook, and Teams at once, demonstrated live with his own Gmail, look into my inbox and tell me the latest email received, then draft a reply to so-and-so. And to a small-business consultant in New York asking whether he could partner with AI experts serving US businesses, the answer was that this is exactly Harshith's own business, building custom dashboards and solutions with partners, and yes.

On Connecting GitHub, Skills, and Grok

Two carry-over confusions were cleared quickly. The skill created in a previous session is not lost: both Claude and ChatGPT carry a built-in skill creator, tag the assistant and type skill to see it, and any skill you have built stays available and is invoked by asking, while the model already knows you through its memory. And GitHub code is not uploaded into Claude or ChatGPT; it is connected. Claude's connector list turned out not to carry a native GitHub option, which became the lesson, because you can still connect anything: say I want to connect my GitHub to you, use GitHub MCP or CLI, and the model inspects the tools already on your machine and finds a route, in his case the GitHub CLI already authenticated with a hundred and forty-three repositories accessible, verified by asking which repository he had worked on most recently, his own website, an hour earlier. CLI, command line interface, got its plain definition, every tool has a terminal version that computers read more easily, and MCP, model context protocol, is simply plugins or connectors, two doors into the same room, with CLI recommended as the easier one. On Grok, asked whether it deserves a place beside Claude and ChatGPT, he was candid that he had tried it once, praised its speed, and explained why with a whiteboard: mixture of experts, where a two-trillion-parameter model routes each prompt to only the small specialist slice that is good at coding or marketing rather than computing across the whole thing, which is why answers arrive fast, though he is no fan of its interface and does his own work in ChatGPT and Codex.

On Multi-Model Tools, and the Council Built Live

The multi-tool question, how to install ChatGPT source code into GetMulti, exposed a misunderstanding worth correcting: GetMulti is a closed product and so are Claude and ChatGPT, there is no source code to install anywhere; it is a council of LLMs, one prompt run through several models side by side, Claude as judge over GPT-5.6, Grok, Mistral, and the rest, useful but token-intensive since every model bills for the same prompt, with the alternative being one powerful model such as GPT-6 Astra in its Ultra mode deploying multiple sub-agents, whose auto-generated names, Einstein and Chandra Shekhar, entertained the room. Then the hour's set piece: because GetMulti cannot look into a codebase, Harshith prompted his agent, mid-conversation, to build a very simple Council of Experts dashboard integrating OpenRouter, the one API that routes to every model, taking inspiration from GetMulti but keeping it simple, then to make it GitHub-connected so it could read his repositories. Four minutes and forty-four seconds later, while other questions were being answered, it was done, a working app with a Strategist on Sonnet, a Challenger, and a Pragmatist on Gemini, his codebase loadable into the council, and one further prompt renamed it Give Multi and published it to a live subdomain of his site. His point to the schools-and-NotebookLM questioner was the same: NotebookLM has no subfolders, one notebook per project, compress files if volume is the issue, and it exposes no API so n8n cannot feed it, but a non-coder just built a custom equivalent while chatting, and that is a product you could sell to schools. The n8n troubleshooting question, an agent node refusing to run for want of a user prompt after a system prompt was added, got the practical answer, open the node's prompt drop-down properly so the entry registers, and save the workflow, because an unsaved workflow does not see its own prompts. The Grok-bots text-to-image-to-video pipeline that was not reaching a hundred percent got his standing advice, build a small front-end dashboard where each agent logs its notes so you can see exactly where the gap is, the way his own newsletter dashboard shows subscribers, opens, and issues, and then fix prompts and integrations against evidence.

On Long-Form Video, and the skills.sh Hack

A request for a quick way to make an eight-to-nine-minute YouTube video met a demonstration from the night before: two videos generated entirely inside ChatGPT, thumbnail included, one a forty-three-second roast of the week's Dario-versus-Trump AI-safety exchange and one a three-minute piece on AI generating enough work to fill the time it saved, both narrated in Harshith's own cloned voice through ElevenLabs, which most of the chat rated close to real. The engine is Hyperframes, a tool that lets an AI agent compose a video by writing code; his plainer description was you type, it creates the video, feeding it images, memes, GIFs, existing footage, and an LLM, with length limited only by the ask and n8n automation the next step once his manual process is proven. The final question, how to train Claude or ChatGPT in a skill you know nothing about, marketing funnels for instance, got the answer the cohort will hear again: the models already hold the skill, what you narrow it with is skills.sh, where a search for marketing funnel surfaced a growth-funnel skill covering acquisition, activation, retention, referral, and revenue, business-school material packaged as a prompt, copied into ChatGPT with install this for yourself, after which help me create a funnel for so-and-so uses it. His closing proof was personal, a self-declared non-coder who hated accounting at school, who built an accounting product on installed skills and sold it to a private equity firm, and a last flourish, a voice skill named Jarvis that, asked live, spoke a thank-you note to the room.

THE MENTAL MODEL

Every tool has a door, an MCP, a CLI, a connector, a documentation link, and you open it by pasting the link into ChatGPT or Claude and saying connect. Confidential work runs locally on n8n plus Ollama, where the computation never leaves your chip. The same models sit under ChatGPT, Copilot, and Azure, so the methodologies transfer. And when a skill is missing, skills.sh installs someone else's expertise, which is how a non-coder builds and sells enterprise software.

Key Takeaways

Stripped to its essentials, the hour made the following points.

- **Plugins are MCPs, and MCPs are a link.** Find the tool's MCP documentation, paste it into ChatGPT or Claude, say connect, and Linear, Notion, or anything else becomes conversational, no code and no engineering team required.
- **Confidential means local.** Install n8n from its public GitHub codebase, run models like Gemma 4 or Qwen through Ollama sized to your machine, and every computation stays inside your own chip.
- **Enterprise tools all have integrations.** Research the platform with ChatGPT, collect its APIs and integrations, and build; Coldbridge, QuickBooks, NetSuite, and Outlook all open the same way.
- **Copilot runs on the same models.** Azure hosts OpenAI, so ChatGPT and Microsoft Copilot share the GPT engine, the course's prompting methods transfer unchanged, and Copilot's only edge is native Microsoft 365 access.

- **Connect GitHub by CLI or MCP.** No native connector needed; ask Claude to connect using the GitHub CLI or MCP and it finds the route itself, after which your entire codebase is a prompt away.
- **Grok is fast because of mixture of experts.** It routes each prompt to a small specialist slice of a two-trillion-parameter model instead of the whole thing, which explains the speed if not the interface.
- **Build the tool you cannot find.** A Council of Experts app on OpenRouter, GitHub-connected and hosted live, was vibe-coded in under five minutes mid-session, and NotebookLM's lack of subfolders and an API is an argument for the same approach.
- **Skills close every gap.** Hyperframes plus a cloned voice makes long-form video from a prompt, and skills.sh turns a growth-funnel or accounting skill into installed expertise, the route by which a non-coder sold software to a private equity firm.

End of Office Hours Two notes. Next: Office Hours Three at 6 pm IST.