

AI GENERALIST ACCELERATOR · 14-DAY AI SPRINT

Day Five

Advanced MCPs: Chaining Tools, Computer Use, and Building Your Own Server

A cleaned, detailed record of the live session and breakout briefing

Topics covered: Why MCPs Exist · Higgsfield Inside Claude · Chaining MCPs for a Launch Film · Computer Use and Routines in ChatGPT · Plugins, Connectors, and the Figma Carousel · Building an MCP Server in n8n · Using AI Versus Using AI Well · The Weekend Ahead

About This Session

Day five returned to the model context protocol introduced on day one and went deep, with a last-minute change of mentor: Dileep, scheduled for the session, had been unwell for three days, so Akshat Bahety, head of the Digital and AI Innovation Center at the Adecco Group and, in Uthappa's introduction, mentor to some five thousand learners, stepped in, with Dileep now due back next week for a session on open-source models and Codex workflows. Uthappa's opening check-in marked the end of the three-day settling window and laid out the weekend ahead before handing over. Akshat's framing was characteristic: if it were easy everyone would do it, problems and learning are not bugs but features of life, and the session is done together, replies in the chat expected throughout. What followed was two and a half hours of live demonstrations, several of which broke on camera and were fixed on camera, which Uthappa later called the day's most valuable lesson, before his own fifteen-minute talk on turning tool knowledge into real implementation and the breakout that built every learner's first MCP server.

HOW TO USE THESE NOTES

The workbook builds a simple email MCP server in n8n, with a walkthrough video at its top, and a separate Additional Workbook under day five resources is for post-session exploration. Because first MCP servers misbehave on every machine differently, Harshith compiled a troubleshooting guide, shared in chat and added to the platform; unresolved errors go to the Ask Doubt channel with a screenshot for a mentor reply within 24 hours. The workbook API keys from earlier days are revoked; make your own free Gemini key or add five dollars to an OpenAI account.

The weekend runs two full-day batches, C14 IST from 10 am to 7 pm IST and C14 EST from 10 am to 7 pm EDT, on both Saturday and Sunday, with no office hours. Your space on the platform tells you which batch you are in, a final switch form was dropped in chat, and whatever you choose stays for the rest of the sprint. Office-hour recordings and transcripts sit under day three on the platform.

Why MCPs Exist

The five-minute refresher started with Uber. Today you open the app, choose a destination, pick a car, and a taxi arrives, three systems, your phone's operating system, Uber, and the physical world, stitched together by you tapping. The wealthy skip the tapping, an assistant is told I want to go to the mall and a car appears, and MCPs exist to do exactly that one thing: take an instruction from you and perform an action on another system. The second example was the OTP on WhatsApp, proof that systems have talked to each other forever through APIs, application programming interfaces, single points of connection that engineers must build, and rebuild, every time a bank wants one more thing from WhatsApp. An MCP is a wrapper on top of that complexity, one connection through which the bank, or you, can send, read, deploy, or request whatever the other system offers, working in both directions and automatically inheriting every new feature the tool ships. His confession about the stakes was personal: he manages his whole life this way, financial accounts, product accounts, work accounts, all in one place, where before MCPs he was clicking seventeen different tools seventeen different times to get data.

Giving Claude a Superpower: Higgsfield Inside Claude

The first demonstration gave Claude something it cannot do on its own, generate images, by connecting Higgsfield, the one-subscription marketplace of image and video models the cohort met on day two. A deliberately lazy prompt for a Japanese town in a hand-painted animation style with animated characters produced, within Claude, a tool call to Higgsfield's generate-image function and an image the chat scored a ten, and the reason it was good was the point: Akshat never wrote the prompt, Claude wrote it and sent it to the tool. Higgsfield's own artifact then offered an animate button inside Claude, a second deliberately bad prompt about schoolchildren walking to school was improved by Claude before dispatch, and a short animated clip with sound came back in three or four minutes, the first video generation triggering a permission request because videos cost more credits. Connecting it took one visit to the tool's MCP documentation for its connector URL and one entry under Claude's settings, customise, connectors, and it prompted the day's recurring vocabulary lesson: ninety-nine percent of connectors are MCPs on the back end, and connectors is simply the friendlier word a product manager chooses for users who never attend a session like this. Two more of his rules landed here. Yes, using Higgsfield through Claude spends Claude tokens and Higgsfield credits both, and that is a very small price against your time; and set a monthly AI budget, five, ten, twenty dollars, and spend all of it, because optimising for cost while learning is how you do yourself a disservice, and experts exist precisely to optimise later.

Chaining MCPs: The Launch Film

The second demonstration chained tools to answer the question the room had voted hardest in the digital world, getting attention, through the use case of a marketing agency that needs to show launch-video work before anyone hires it. Firecrawl was introduced first, a scraper that extracts information from any website and whose job, not yours, is to get past whatever blocks scraping. Then one prompt to Claude: use Firecrawl to fetch Nike's brand guidelines for video content, then pitch and generate a realistic launch video, people running in shoes, with Higgsfield. Claude researched the guidelines, wrote shot-level pitches, respected a rule it found about logo use, generated hero images the chat rated highly, and produced multiple short clips because it worked out on its own that the models cap at three to five seconds. The honest wobble came at the end: asked to merge the clips, Claude in chat could not save files, so Akshat downloaded them himself and attached them, with the note that in cowork mode, where Claude has file access, it would have done it, and the line that followed became a refrain, if Claude cannot do something behind a screen, it is a skill issue, and he will figure out a way. The stitched ad, roughly twenty seconds with audio, was rated an eight to ten by the room and by him, and the outcome he insisted on was not the video's polish but that it came out of a single chat without learning Higgsfield, Firecrawl, or brand guidelines, work that would once have cost a full day of a skilled person's time.

Computer Use and Routines in ChatGPT

While the clips rendered, the session switched to ChatGPT's desktop app to show a feature without any connector at all: the built-in browser in work mode, a toggle panel that opens a real browser inside ChatGPT where you can log into any site, in this case Twitter, logged in before the session so no password went on screen. The prompt was small, go through my feed, pick one interesting tweet, draft a reply, and the demonstration was Akshat's hands held in front of the camera while ChatGPT scrolled the feed by itself, checked the layout, chose a tweet, wrote the reply, and, on a second instruction, posted it. The skill is called computer use, it works on any website that opens in Chrome, LinkedIn, Gmail, anything, and it was named the current gap between the two tools, Claude does not yet do this as well. The follow-up automated it: a single sentence asking for a scheduled routine that reviews the feed daily and sends five high-quality opportunities to reply to produced a routine, prompt and all, set for 9 am and then moved to the evening because nobody writes good tweets at 9 am, visible under the Scheduled tab alongside the India-headlines routine he ran every morning for a year. His advice on the mechanics was to ignore them, do not decide whether ChatGPT should use Codex automation or anything else, state the outcome and let it choose, because the playbook of the fourteen days is to become the manager who runs agents, the person whose ten thousand ideas no longer die for lack of skills, time, or money to hire.

Plugins, Connectors, and the Figma Carousel

The naming confusion got its definitive treatment: a Gmail plugin in ChatGPT is a connector in Claude is an MCP underneath, plugins equal MCPs equal connectors, and no company will agree on one word because, in his phrase, marketing has destroyed more products than engineering ever will, Salesforce's serial rebrands being exhibit A. Your job is not to become an expert in what a word means in a given tool but an expert in using and building with it, and the trend behind all of it is that every tool is building an MCP from day one because the next generation of users is not only humans but AI agents, who like MCPs precisely because everything travels as text. His own purchase rule follows: software without an MCP does not get bought, because learning UIs is now the AI's job, not his. Financial MCPs exist too, Robinhood in the US and Zerodha in India, and trading through an agent is possible, offered explicitly as information and not advice. The demonstration that made the point cost nothing extra: with Figma connected as a ChatGPT plugin, one prompt asked for a four-slide Instagram carousel explaining photosynthesis in an easy, animated way, and ChatGPT built the designs directly inside a Figma file, editable element by element, exportable as images for LinkedIn or Twitter, and paid for only in ChatGPT tokens since Figma's free plan carries the design work. The refrain closed the section: tools will keep changing, your job is to choose the task and the outcome, use every tool you have, and if a company ever pays you to use theirs, use that one.

Building Your Own MCP Server in n8n

The centrepiece was building an MCP server from scratch, with Perplexity search as the example and a warning repeated three times that it is an example, chosen for the simplest widest use case, not a recommendation about Perplexity. The whiteboard version came first: everything in MCP is a URL, behind every MCP URL sit tools, tools can be anything, a search, an image generator, a calculator, and on the back end each tool is just an API, so an MCP is a wrapper around tools that are wrappers around APIs, and your AI only needs to know about the outer layer. The n8n version mapped onto it step by step. Search for MCP and add the MCP Server Trigger, which immediately supplies the MCP URL. Add a tool by searching for HTTP Request, because HTTP is how the entire web works and every organisational API is reached through it. Configure the tool from Perplexity's own documentation, a Google search away, by copying the curl command for the search API, the sentence he asked the room to remember being if you want to use our API, use our curl command, and pasting it into n8n's import-curl box so nobody has to learn methods or parameters. Add an API key from the Perplexity platform, never on screen. Rename the tool Perplexity Search, publish, copy the production URL, and in Claude add it under connectors as a new MCP URL, no sign-in required. Then reality arrived. Claude reported the connector broken, and the first fix was a body parameter in the HTTP tool that needed switching to let AI decide; the second failure took a full minute of live debugging before the cause surfaced, the API key had not been saved after the edit; the third attempt, a search for Sanju Samson, came back with Perplexity's results inside Claude, to a chat that had watched every step. The lesson was explicit, things break, and your job is to fix them with AI, and the restaurant analogy closed it: you order, the kitchen does steps two, three, and four, and you never visit the kitchen because you are there to eat, not to cook. Uthappa added the why: Claude's connector list holds a few hundred servers at most, Whisper Flow for one exposes an MCP URL without a native listing, and any tool with an API endpoint, or one you vibe-code yourself, becomes usable inside Claude, ChatGPT, or Codex through exactly this non-technical route.

Akshat's Own Use, and the Learning Rule

The wrap was personal. His entire financial planning now runs through MCPs, email, portfolio, statements all connected to Claude, which tells him monthly what needs doing, a task that once took a couple of days and now takes two hours. On security and privacy he was blunt and unbothered, nothing is fully private anywhere, the tools carry OAuth and proper authentication, and the useful response is to build anyway. The challenge to the room was one use case each, and the simplest one revealed the gap, fewer than a fifth of the chat let AI read their email daily. Buy back time in proportion to what you spend, twenty, forty, sixty dollars a month should return that much in hours, and learn the way he learned to never forget an API key again: you are not learning when everything works and it is fun, you are learning at midnight banging your head on the table asking why this does not work.

Using AI Versus Using AI Well

Uthappa's fifteen minutes began with a slide of a knotted brain and a poll that split the room, most still figuring out how to make any of this work at their jobs, and his reassurance was that this is the normal starting point, that it took him a year, and that there is no shortcut: skills first, then thinking in workflows, then designing systems, the same arc as the computer shrinking from a box in the office to a phone in a pocket. His own mistake was building imaginary workflows and endless vibe-coded products for a future self instead of the work he actually did, and the fix was observing his day for where time leaks, handoffs stall, and effort repeats. The system that followed, twelve months of building and rebuilding, was shown as a diagram: Granola, Whisper Flow, and Obsidian feed meetings, thoughts, and research as raw input into a local folder; Claude Code, which the cohort meets next week, organises it into work and personal knowledge that becomes the shared memory; Supabase holds tasks captured by voice and visible to the team; Hermes and Antigravity ship content from that context; and Codex absorbs Slack and email into the tracker so nothing needs reading by hand, all wired through MCPs. It is not a template, it changes per person, and it takes six to twelve months to reach the point of working with AI as a partner. For companies he borrowed Reid Hoffman's observation that most enterprise AI strategies run backwards, chief AI officer, pilot, showcase, while the value sits in boring work nobody writes a press release about, the forty-hour weekly report cut to twenty minutes, fifty people each taking back four hours, and the company building a memory of what AI is actually good for. The cautionary tale was a client whose leader bought Claude licences for everyone while IT never cleared Outlook to connect, so everyone had a licence and nobody had a use case; the instruction was to learn now and connect later, because the approval will come. The AI champion every B2B client asks for is spotted by one test, does this person know more than ninety percent of the people around them, which after five days most of the room does, a gap with a shelf life to be spent on one workflow at work this week.

The rest was guardrails for the ambitious. On whether AI takes jobs, Vaibhav's honest I don't know, plus a Berkeley study in Harvard Business Review of forty engineers who adopted AI voluntarily and got no free afternoons, only wider scope, more code review, more coaching, and burnout, so his prescription is two or three tasks a week done without AI, writing the email, doing the analysis, to keep alive the skill you will need on the day a confident output is wrong. On ambition, anyone using AI to cut costs will be beaten by the company next door using it to serve more customers, the arithmetic being not the same crore with five people but five to ten crore with more, which is why Outskill itself is hiring. On automation, ninety percent is the limit for anything touching a client, a payment, or a person's record, illustrated by the invoice workflow they refused to fully automate, stopping at a draft in the

outbox so a human checks tone, amount, and attachment before send, because agents chasing invoices while you sleep is content, not operations. On tools, Fei-Fei Li's line from the Huberman Lab podcast, start with the problem not the tool, and stop asking whether Claude beats Codex this month, both will solve it, pick the one that resonates. On decisions, use AI to widen the options, competing hypotheses, three strategies you did not consider, the strongest argument against your plan, while the decision stays yours, especially the one about who stays and who goes, where the agent gives context and nobody gets to say the agent decided. Two closing images: the person who talked to AI ten minutes every morning for three months until the answers sharpened on three months of context, and do not use AI the way a drunk uses a streetlight, for support rather than light. And the essay he asked everyone to search, Paul David's The Dynamo and the Computer, on the forty years factory productivity did not move after electrification because owners swapped the steam engine and kept the driveshaft layout, until someone put a small motor on every machine and redesigned the floor; we are at the shaft-and-belt stage of AI, and the better question than where do I insert AI into my process is what work is now worth doing that was not worth doing before.

THE MENTAL MODEL

MCPs take an instruction from you and act on another system; connectors, plugins, and MCPs are one thing under three marketing names, and any tool with an API can become one through an n8n MCP Server Trigger, an HTTP Request tool, and an imported curl command. Chain them inside one chat and let the AI write the prompts, choose the sub-tools, and fix what breaks. Then use AI well: pick one real workflow this week, automate to ninety percent with a human on the last ten, keep the decisions, and ask what work is now worth doing that was not before.

The Breakout, and the Weekend Ahead

The workbook put the lesson into every learner's hands: build a simple email MCP server in n8n, not because one is missing but so the mechanism sticks, with Uthappa's own future use in mind, a server that reads five inboxes at once where the native Gmail connector allows only one. Ninety minutes of rooms produced the expected spread, first servers that worked and first servers that did not, which is why Harshith's troubleshooting guide went out in chat and onto the platform and why the Ask Doubt channel takes screenshots for a mentor reply within a day. Two housekeeping answers stood out: the infographic of Uthappa's system was drawn by asking Claude, which cannot generate images without an MCP, for a prompt that then went to ChatGPT; and the bonus tools the cohort holds are deliberately not the only tools taught, with deep dives on each living in the content library. The weekend is the sprint's turning point. Saturday's first half teaches voice agents and its second half building products with AI, vibe coding, completing all five levels of the curriculum. Sunday opens with Vaibhav, who is flying to the US for MetaConnect immediately after, on how to connect everything learned so far and how people outside are building and monetising with the same skill set, and closes with the cohort building version one of Jerry, its own multi-agent assistant, in n8n. Full days, 10 am to 7 pm in your batch's zone, food ready, rest tonight.

Key Takeaways

Stripped to its essentials, the session made the following points.

- **An MCP takes your instruction and acts on another system.** It is a wrapper over the APIs systems have always used to talk, working both ways and inheriting every new feature a tool ships, which is why one chat can replace seventeen tools clicked seventeen times.
- **Let the AI write the prompts and pick the tools.** Higgsfield inside Claude produced a ten-out-of-ten image and an animated clip from two lazy prompts because Claude improved them before dispatch; your job is orchestration, not prompting every tool by hand.
- **Chain tools in one chat.** Firecrawl fetched brand guidelines and Higgsfield built a launch film in a single conversation without learning either tool, and when Claude could not download files in chat, the answer was cowork mode and a refusal to call it anything but a skill issue.
- **Computer use needs no connector.** ChatGPT's built-in browser scrolled Twitter, drafted, and posted a reply on its own, and one sentence turned that into a daily scheduled routine; state the outcome and never decide the mechanics for it.
- **Plugins, connectors, and MCPs are one thing.** Every tool is building one because agents are the next users, Figma via plugin designed an editable carousel for tokens alone, and software without an MCP is software Akshat does not buy.
- **Any API becomes an MCP server in n8n.** MCP Server Trigger for the URL, HTTP Request for the tool, the vendor's curl command imported, an API key, publish, paste the URL into Claude's connectors; it broke twice live, a body-parameter toggle and an unsaved key, and worked on the third try, which was the lesson.
- **Use AI well, not just use AI.** Skills become workflows become systems over six to twelve months; pick one real workflow this week, spend the ninety-percent knowledge gap while it lasts, automate to ninety percent with a human on the last ten, do a few tasks a week unaided, and keep every decision about people your own.
- **Ask the dynamo question.** Factories gained nothing from electricity until they redesigned the floor around small motors; the better question than where to insert AI into the old process is what work is now worth doing that was not worth doing before.

End of Day Five notes. Next: voice agents and vibe coding — Day Six.