

AI GENERALIST ACCELERATOR · 14-DAY AI SPRINT

Bonus Office Hours Four

Office Hours Four: APIs, Webhooks, and MCPs, and Taking Jerry Into the World

A cleaned, detailed record of the first week-two Slido question hour

Topics covered: API Versus Webhook Versus MCP · Jarvis: Jerry on MCP Servers · Voice and WhatsApp Inputs · Permissions, Sharing, and Guardrails · Building for Employees and Clients · Tools for Content and Websites · Cybersecurity and Agent Breaker · Local Hardware and RunPod · Debugging With n8n's Assistant

About This Session

Week two's office hours opened at 6 pm IST with a new mentor answering, Samruddhi Mokul, founder of Pace Labs and a week-one breakout regular, and Shasti hosting a room she greeted as survivors of the two full weekend days, with the reminder that the hour runs daily through day thirteen and that anything unanswered rolls into the next. The Slido format held, questions upvoted into order, and the mentor drew her standing boundary early: programme, resource, and bonus-tool questions belong with Shasti and Sumedha through support and the Ask Doubt channel, not the mentor. One housekeeping note landed before the first question, the workbook API keys had expired at 6 pm IST, minutes earlier, and were revoked. The hour was shaped by the Jerry build the cohort had finished the night before, most questions being some version of how to make it real, and by Samruddhi's own version of the same idea, an assistant she calls Jarvis, whose workflows she shared as importable JSON along with a promise to compile every tool, link, and community she mentioned into one document for the platform. Her connection dropped twice to tree-cutting on the wiring outside, and she carried on.

HOW TO USE THESE NOTES

Three JSON workflows were dropped in chat and go to the office-hours resources on the platform: the Jarvis master orchestrator with its MCP servers, a Telegram trigger that accepts text and audio, and a WhatsApp customer-support flow that handles voice, image, and text. Import any of them from n8n's three-dot menu, Import from File, add your own credentials where the red exclamation marks appear, and they run.

Local installation of n8n and its desktop setup was deliberately not demonstrated, because it takes a full session; the platform carries a detailed video and documentation for it. Office hours continue tomorrow at 6 pm IST.

API Versus Webhook Versus MCP, Three Ways

The question that recurred all hour, when to use which, got the phone-call framing first. An API lets two pieces of software talk, but only when you make the call: you ring the other software, it picks up, it answers, and if you want to know whether a job is done you must keep calling. A webhook is the reverse, you are already on the line, you have asked for something, and when the work is done the other side comes back and tells you, so webhooks sit live and fire only when a specific action happens at the other end. An MCP is a newer layer on top of APIs, built for AI agents: an MCP server comes connected to every API endpoint a tool exposes, so instead of handing an agent a hundred tools and watching it hallucinate over which to pick, you hand it a few MCP servers and let each route to the right call. The second telling, prompted by a learner who had imported a webhook in the breakout without building one, added the n8n specifics: every workflow can start with a webhook node, which issues the URL; when another website uses it you switch to the production URL and keep the workflow live; and whether the workflow ends with a Respond to Webhook node depends on whether anything needs to come back, yes for a customer-support agent that must answer the chat, no if the job is only logging rows into Google Sheets or a database. The third telling answered can you show API and MCP the way you showed webhooks: yesterday's Fireflies HTTP request was an API call read straight from the vendor's documentation, one action per endpoint, one endpoint for transcript IDs and another for the transcript of a given ID, and every n8n node that talks to a service is itself an API call dressed as a node. MCP, in n8n, is an MCP Client tool added from the agent's Tools plus button with the server URL inside, because n8n has no connectors panel the way ChatGPT and Claude do.

Jarvis: Jerry Rebuilt on MCP Servers

Her Jarvis made the MCP case visible. It is Jerry with more hands, task manager, calendar, meetings, finance, and Google Contacts, plus a thinking tool that complex tasks are routed to for a more considered answer, taking text and audio from Telegram and replying there when a task is done. The architecture is six separate n8n workflows talking to one another: a master orchestrator and five MCP servers, one per domain, each holding the tools that yesterday were attached directly to sub-agents. The advantages she listed are the reasons to graduate from Jerry's design when a system grows. Fewer, cleaner choices for the master, so less hallucination over which tool to call. Reuse, because a calendar MCP server built once can be plugged by URL into any other workflow, Jerry included, instead of re-adding ten calendar tools every time. And security, since for a client with hundreds of connections you build and secure MCP servers per client and connect only those to the agent. The production-readiness question, how to turn Jerry into a reusable system for small-business clients with agent roles, secure access, human approvals, shared memory, and failure handling, got the same shape: everything can run in n8n, but anything stored goes to the client's own backend, Supabase, AWS, whatever they run, through the nodes already there; roles come from the flowchart drawn first, without which nothing is known; and approvals use n8n's human-in-the-loop node, routing data to a named person before the workflow proceeds.

Voice, WhatsApp, and Jerry in Daily Life

Several learners had expected Jerry to answer a voice message rather than a typed one, and the answer was a shared workflow rather than a lecture: the trigger must accept audio, Telegram does, and then one extra step decides what arrived, text goes straight to the master orchestrator, audio is transcribed to text first, because agents can process audio directly but it is expensive. The WhatsApp version she shared handles three inputs, voice transcribed, images analysed, text passed through, all converging on the master agent. On the limits, Meta's official WhatsApp APIs allow only business accounts as triggers or as MCP tools, so personal WhatsApp routes through Unipile, a service that exposes API keys for actions on a personal account and can itself be wrapped as an MCP server. How to use Jerry every day, then, is simply a trigger question: swap the chat trigger for whatever you already live in, Telegram, WhatsApp, email, add several triggers at once if you want to reach it from several places, or start with a webhook trigger when the front door is a website. The running-locally questions got the same practical edge, a local n8n runs only while the laptop is on, so host on Railway if it must run unattended, and downloading n8n locally does not mean downloading a model, the authentications and connections work exactly as before. A downloaded ChatGPT app is still chargeable; only genuinely open-source models are free.

Permissions, Sharing, Guardrails, and Building for Others

A worried question claimed that unauthenticated Gmail nodes leave every inbox open to anyone with the production link, and the correction was precise: when you signed in with Google and created credentials you granted permissions, and permissions do nothing until the workflow is triggered or executed, so give all the permissions the pop-up asks for, or the job fails when it runs. Sharing a workflow with a friend is safe for the same reason, download the JSON and send it, credentials are encrypted and never travel, the recipient sees red exclamation marks until they add their own, and n8n's project workspaces keep credentials from leaking even between your own projects. Publishing does not expose a workflow to others either; n8n's built-in publish checklist applies its own guardrails, and for chat triggers she showed her own filter node that passes messages only from her user ID. Building Jerry for a company's employees raised the authentication question properly, and the answer is never to collect other people's OAuth sign-ins: every platform has an admin control, Google's Cloud Console for one, where you create API keys with configured scopes, five functions for one employee, three for another, the same route for clients. On cybersecurity for vibe-coded products, two habits: always add a verifier layer, a step where AI double-checks its own work, because LLMs say yes to everything, and run the codebase through QA and security tools built for the purpose rather than hand-checking it. Her recommended way to understand the attacks themselves was a game, Agent Breaker by Lakera, which she called one of the best companies in AI-agent security, where you hack your own AI level by level and learn prompt injection by doing it.

Tools for Content, Websites, and Ideas

The what-should-I-use question for a personal AI-news channel got a short, opinionated list, with the promise of a fuller document later. Image and video generation, Higgsfield, the best at this point for both; cloning yourself, HeyGen; automated posting and personal branding, Blotato, spelled out for the chat. Website building was tiered by backend weight: Lovable for a fancy, beautiful front end with a light backend, Replit or Bolt when a real backend is needed, and Cursor, Claude Code, or Antigravity when the backend is heavy. On the full-stack companies Vaibhav had described and how to start, she reframed the question: every product on the market solves a specific problem, and a beautiful product for a problem nobody has will not sell, so find the problem statement first, then build, with RAG a legitimate product among many. For finding problems and demand her two go-to sources were Product Hunt, where new products launch and trending categories reveal what people want, and Reddit's communities, with a promise to list the ones she uses. The Fireflies API location, asked about for steps seven to nine of the workbook, was two clicks, Settings, then MCP and API, where both the API key and Fireflies' own MCP server live.

Hardware, Local Models, and Debugging

Laptop requirements for local models depend on the model, and the runners say so: Ollama, LM Studio, and Msty label each model green, yellow, or red for your machine, so check the label and run what is green; a sixteen-gigabyte laptop is fine for an agent doing video creation. When heavier work arrives there is no need to buy hardware, which is her own practice even on an M5 Pro, because RunPod rents GPUs by the hour, pay for one hour and use one hour, the way she trained models on terabytes of data. Offline RAG runs inside LM Studio or Msty, which have knowledge-base features that accept your documents locally, while a RAG meant for a website is built on a vector store like Pinecone and queried from workflows. The Retell error a learner hit, HTTP 400 with an email-validation message, was Cal.com rejecting a badly formed email or phone number, fixed by prompting the agent to collect them in the right format and diagnosed by copying the whole error JSON from Retell's tool-result dropdown into Claude, which is how the mentors debug in the rooms too. The context-window error, input exceeds the model's window, was the day-seven lesson returning, and for every n8n error her advice was the platform's own AI assistant, which reads the workflow, spots the failing node, explains the error, and, told to resolve it, changes the settings on your behalf. Smaller answers filled the gaps: n8n autosaves, the save button in the workbook belonging to older versions, and the admin panel's Manage view shows your version and lets you upgrade to the latest stable release; Google Sheets columns appear automatically once the sheet is attached, and values are dragged from the input panel, which shows every previous node's output; Retell agents cannot be reached by anyone until a phone number is attached, so unpublishing is rarely needed; and a plain n8n workflow with no AI node is just automation, if-this-then-that, whereas Hermes and its kind are agentic systems that learn your actions and preferences over time, most of them running on MCP servers underneath.

THE MENTAL MODEL

You call an API, a webhook calls you, and an MCP server routes an agent to the right API without a hundred tools to choose from. Permissions do nothing until a workflow runs, credentials never travel with a JSON, and other people's access comes from admin-console API keys, not their logins. Swap Jerry's trigger for the app you live in, transcribe audio before the orchestrator, verify AI's work with a second layer, and debug by handing the whole error to Claude or to n8n's own assistant.

Key Takeaways

Stripped to its essentials, the hour made the following points.

- **Call, be called, or route.** An API answers only when you call it, a webhook fires when something happens at the other end, and an MCP server is a layer over a tool's APIs that keeps agents from hallucinating across a hundred tools.
- **Jarvis is where Jerry grows.** One master orchestrator plus one MCP server per domain, each reusable by URL in any workflow and securable per client, is the architecture for anything beyond a prototype.
- **Voice is one transcription step away.** Telegram or WhatsApp triggers accept audio and images; transcribe or analyse before the orchestrator, and route personal WhatsApp through Unipile since Meta's APIs are business-only.
- **Permissions are not actions.** Granting Google scopes does nothing until a workflow runs, shared JSONs carry no credentials, publishing exposes nothing, and employee or client access comes from admin-console API keys with scoped functions.
- **Verify, then test with tools.** Add a verifier layer to every vibe-coded product, run QA and security tools on the codebase, and learn the attacks by playing Lakera's Agent Breaker.
- **Rent, don't buy.** Run green-labelled models locally on Ollama or LM Studio, and for anything heavier rent GPUs by the hour on RunPod; host unattended workflows on Railway rather than a laptop that sleeps.
- **Find the problem before the product.** Product Hunt and Reddit reveal demand; Higgsfield, HeyGen, Blotato, and a Lovable-to-Cursor ladder cover content and websites by backend weight.
- **Let the assistants debug.** Copy the whole error JSON into Claude, or ask n8n's built-in AI, which reads the workflow, names the failing node, and fixes settings when told to.

End of Office Hours Four notes. Next: Office Hours Five at 6 pm IST.