

AI GENERALIST ACCELERATOR · 14-DAY AI SPRINT

Bonus Office Hours Five

Office Hours Five: OpenRouter, Error Workflows, and the Limits of Autonomy

A cleaned, detailed record of the fifth Slido question hour

Topics covered: WhatsApp, Forms, and Triggers · Local n8n and Its Costs · Code Nodes Before Agents · OpenRouter and FAL · Domains, GitHub, and Respond to Webhook · Three Tiers of RAG · Error Workflows · Bots, Agents, and Autonomy · The Research-to-Video Pipeline

About This Session

The fifth office hour ran at 6 pm IST with Sumedha hosting and Samruddhi Mokal answering for the second day running, starting with the questions carried over from the day before and taking the rest by Slido upvote. Her boundary held as always, programme questions, recording delays, and where resources sit on the platform were parked for Sumedha, who answered them at the close: Zoom takes longer to return the files of long sessions, the team QC's them before upload, and Zoom's policy forbids recording breakout rooms, which is why workbook demo videos exist instead. Sumedha also set expectations on the workflows Samruddhi shows, she charges clients for them and cannot share every one, so learn to build them and ask when stuck. The hour's centre of gravity was production practice, how to keep a client workflow alive through a provider outage, how to structure a RAG for its audience, and a candid reality check on how autonomous today's agents actually are.

HOW TO USE THESE NOTES

Two shared resources are worth finding in the office-hours section of the platform: the half-hour video Samruddhi promised on connecting Gmail to a locally hosted n8n through Google Cloud Console, and the shape of her error-notification workflow, an error trigger feeding an email or Slack alert, which she sets up for every client workflow.

Before the course ends, open n8n's Credentials tab, and the equivalent in every app you authorised, to review and revoke what you no longer need. Questions still in the queue roll to tomorrow's hour; the main session with Dileep starts sharp at 7.35 pm IST.

WhatsApp, Forms, Triggers, and Running Locally

The WhatsApp question got the same two-lane answer as the day before, sharpened: n8n has a WhatsApp node, and through Meta's official APIs it connects only business accounts, so a personal number goes through Unipile, whose API keys expose actions on a personal account. A follow-up asked whether the tool blocks accounts, and her answer was a story rather than a policy, she has never seen Unipile cause a block, but WhatsApp itself blocked her for twenty-four hours after she blasted a family-function invitation to her whole contact list, no AI involved, so stay within the platform's limits and, when AI sends, personalise each message so the wording differs even when the context is the same. Two trigger misconceptions were cleared: Google Forms and Google Sheets are separate tools with their own triggers, so a form submission or a new sheet row starts a workflow through its native node, and webhooks are only for cases where no node exists. Antigravity's two downloads were untangled, the EXE is the coding agent, the IDE is the editor environment like VS Code with the agent inside it, .exe is just a file format, and the IDE is not tied to a browser. Running n8n locally got its honest cost-benefit: free, but workflows run only while the laptop is on, the cloud version's AI assistant is missing so bugs must be solved alone, and every connection is technical, Gmail in particular requiring a web app in Google Cloud Console with each of Gmail's separate API scopes ticked and a client secret pasted into the node, a half-hour process she promised as a video. Connecting a local n8n to Claude as an MCP needs a public URL, so a localhost address must be tunnelled through ngrok, covered in the platform's local-hosting video; and n8n currently has no MCP server for ChatGPT at all, only for Claude.

Code Nodes Before Agents, and Tokens

Two questions about messy inbound data, B2B orders arriving over email and WhatsApp for an ERP, and unstructured email data for a CRM, got the same discipline: check whether the ERP already has a data-cleaning layer, and if not, add a cleaning step in n8n, but not an AI agent. If the transformation is keyword-to-keyword or purely a format change, a Code node running a little JavaScript or Python does it, and n8n's AI assistant will write that code from a description of the input and the wanted output; an agent is reserved for work that needs dynamic thinking, because adding one where none is needed only adds cost. The large-PDF token question had no hack, only choices: paste just the part you need, upload the whole file if it is a one-time task, and build a RAG if it recurs, which saves a great deal over time. Bot versus agent was defined, a bot is if-else branching like the old WhatsApp menus, an agent has its own reasoning and tools and acts dynamically, though industry uses the words interchangeably. And an OpenAI key that misbehaved after the workbook keys expired was traced not to mismatched email addresses, which are irrelevant, but to provider load, since OpenAI and the others prioritise heavy-usage tiers when demand spikes, so reconnect or create a fresh key.

OpenRouter, FAL, and Keeping Workflows Alive

OpenRouter was demonstrated and defended twice. It exposes hundreds of models behind one API key, so the first reason is money, five dollars of credit there instead of five dollars each at OpenAI, Anthropic, and Gemini just to test; the second is uptime, shown by opening Anthropic's status page to a recent three-hour partial outage across Claude, the console, and the API, because when a provider goes down you cannot tell a client the software is wrong, and with OpenRouter you switch models inside the same node without touching credentials or configuration. She opened Jerry, swapped its OpenAI chat model for the OpenRouter chat model, and scrolled the list, adding the client rule, one OpenRouter key per client so usage stays trackable. The models are identical to the vendors' own, GPT-5 Mini on OpenRouter is GPT-5 Mini, nothing selected or altered. The threshold question, when to stop paying individually, answered itself from those two reasons. OpenRouter is for LLMs, text models; for image, video, and audio generation the equivalent one-key marketplace is FAL, which is what she recommends for video production. For genuine resilience there are two mechanisms: an If node that reroutes when a model call fails, and, as best practice for every client build, an error workflow, a separate workflow starting from n8n's error trigger and assigned to the workflows it guards, which fires on any failure, LLM-related or not, and sends the details to email, Slack, or Teams so a tech team can act.

Domains, GitHub, and Respond to Webhook

Connecting a vibe-coded site to your own domain was explained by analogy to API keys: Lovable or Bolt's Publish offers a custom-domain option, you buy a domain from a registrar, and you paste its DNS record into that box the way you paste a key; a fully custom-coded site from Cursor or Claude Code first goes to a host like Vercel, then connects a domain the same way. GitHub, asked about by someone who missed day eight, got its Google-Drive-for-code summary, a main branch that is what users see live, feature branches per person that merge only after the repository owner's review, and one command, git pull, that fetches an entire codebase instead of file-by-file downloads, which is also the escape hatch from choosing the wrong tool: push Lovable's code to GitHub, pull it into Claude Code, and keep building on top. The Respond to Webhook confusion from the day-eight workbook was resolved in two prompts: the node passes on whatever its previous node outputs, so define that output, then in Lovable's same chat first configure the webhook with its URL and second ask for the complete response of the webhook including the fields you want shown on screen. A publish button greyed out after pasting someone else's workflow simply means unresolved red exclamation marks, credentials that never travel with a JSON, fixed by adding your own or by asking n8n's assistant to list the errors.

Three Tiers of RAG, and How Autonomous Agents Really Are

How to create a RAG depends on who uses it, so she gave three tiers. For yourself alone, a project in ChatGPT or Claude is already a RAG, upload the documents and it answers from them, the simplest path. For a limited group, a custom GPT with its own knowledge base, shared by link to you only, to people with the link, or publicly, still without technical setup. For chatbots serving many users or a website, a vector store such as Pinecone or Supabase, with an agent retrieving from it, the day-four architecture. The desktop assistant a learner remembered from an earlier session, listening and answering as if a person sat beside the mentor, was demystified as browser use or desktop use, already present in Claude Desktop and Codex, which watch the screen's pixels and act, nothing extraordinary once you find the use case, and the only route into on-premise software with no API, since without APIs n8n cannot connect at all. The question that produced her longest answer challenged the course itself, whether spending every day on n8n is wasted when the industry is moving to agents that plan, reason, and execute. Her reply was that those capabilities exist with real limits: the decades-old idea of an AI forward-deployed engineer that meets clients, gathers requirements, plans, and builds the automations is not possible today, first because agents need data and most companies' data is uncleaned and unstructured, so agents hallucinate on it, and second because the technology is heading there but is not there, live avatars talking on your behalf being one visible constraint since video generation is slow; what looks like it in marketing hides the challenges you meet when you go deep. Related definitions closed the thread: Antigravity generates code from prompts while n8n builds workflows, Lovable is also a vibe-coding tool, and the difference between a normal agent and an autonomous one is the trigger, one runs when you invoke it, the other is always on and decides for itself when to act.

Tools, Scrapers, and the Research-to-Video Pipeline

The tool questions were answered in her usual short form. Voice cloning, ElevenLabs; video cloning, HeyGen; research, Perplexity, and when research means scraping social media, Apify, whose marketplace calls its scrapers actors, chosen by the most ratings and the most users, with the warning to prefer scrapers that do not store cookies because those risk your account, and BrowserAct, an n8n for scrapers, when a scraper must take actions like downloading. Images, Higgsfield or FAL, never Canva, which is for decks, and never Lovable, which builds websites; Adobe's own AI is fine for creatives but not for websites, where Lovable wins on front end. Websites with CMS capability are a backend question, Lovable for simple, Bolt or Replit for moderate, Antigravity, Claude Code, or Cursor for heavy, all connecting Supabase so the tables and authentication are created for you. A Microsoft-only company's n8n equivalent is Power Automate, node-based and built for the Microsoft suite; Slack, for anyone who asked, is simply Teams for organisations that do not use Teams; and Bitbucket connects to Cursor through code, extensions, or the terminal, in that order. The last question, how to automate research into cloned videos into posts, got a walk through her own marketing workflow: a research agent files content ideas into a Google Sheet; a video generation branch pulls each row, writes the title and the prompts, generates image, video, and audio, then merges the three in a rendering platform, Creatomate in her case, to produce a final video link for publishing; a parallel branch writes post text and generates images for social; and a clone branch routes through HeyGen so voice and face are hers. Rendering, she noted, is the step most people forget to plan for.

THE MENTAL MODEL

Use native triggers before webhooks, code nodes before agents, and projects before vector stores. One OpenRouter key buys every text model and instant failover; FAL does the same for images and video. Guard every client workflow with an error workflow, keep code in GitHub so you can change tools without starting over, and remember that fully autonomous agents remain a marketing story until the data underneath is clean.

Key Takeaways

Stripped to its essentials, the hour made the following points.

- **Native nodes first, webhooks second.** Google Forms and Sheets have their own triggers, WhatsApp has a node for business accounts and Unipile for personal ones, and webhooks are for what no node covers; stay within WhatsApp's send limits and personalise each message.
- **Local n8n is free and costly.** It sleeps with the laptop, has no AI assistant, and connects to Gmail only through a half-hour Google Cloud Console setup; expose it to Claude with an ngrok tunnel, and host on a server for anything that must run unattended.
- **Code nodes before agents.** Format changes and keyword filtering belong in a JavaScript or Python code node the assistant can write for you; agents are for dynamic thinking, and anything else is cost without benefit.
- **One key, every model, instant failover.** OpenRouter saves the five-dollar minimums and lets you switch models inside the same node when a provider goes down; FAL is the same marketplace for creative models.
- **Every client workflow gets an error workflow.** An error trigger assigned to the workflows it guards fires on any failure and alerts email, Slack, or Teams, with an If node as the in-line fallback.
- **RAG in three tiers.** Projects for yourself, custom GPTs for a group, a Pinecone or Supabase vector store with a retrieving agent for many users or a website.
- **Autonomy is not here yet.** Agents plan, reason, and act within limits, uncleaned company data makes them hallucinate, and the forward-deployed AI engineer of the marketing decks stays a story for now, so n8n time is not wasted.
- **Pick tools by backend weight, keep code portable.** Lovable, then Bolt or Replit, then Antigravity, Claude Code, or Cursor, all on Supabase; push to GitHub and pull into the next tool when requirements grow, and plan for rendering when pipelines end in video.

End of Office Hours Five notes. Next: Office Hours Six at 6 pm IST.