

AI GENERALIST ACCELERATOR · 14-DAY AI SPRINT

Bonus Office Hours Six

Office Hours Six: Going Live, Getting Paid, and Guarding the Agents

A cleaned, detailed record of the sixth Slido question hour

Topics covered: Designs and Agents on a Live Website · Templates, Voice Inputs, and Custom GPTs · Claude Code Needs No Harness · Licences and Privacy · Zapier Versus n8n · FAL Pricing and Video Pipelines · Orchestrator Architecture · Three Ways to Get Paid · API, Webhook, MCP · Hybrid RAG, Guardrails, and Finding Clients

About This Session

The sixth office hour ran at 6 pm IST on day ten with Sumedha hosting and Samruddhi Mokal answering for the third day running, Slido open and questions taken by upvote as usual, with programme questions parked for Sumedha's close. The hour had a distinctly professional bent: how to put a design or an agent on a website that already exists, how to structure and charge for automation work, and how to keep agents from doing what they should not. Samruddhi confirmed the resource guide she is compiling, every tool, repository, and workflow mentioned across her hours, including the local-n8n Gmail video and her own toolkit marked as favourites, targeted for the end of the week, and offered the video-generation template from her marketing workflow while declining, understandably, to share the whole thing. Sumedha's wrap pointed the room at the recordings and transcripts under each day's resources, the official WhatsApp groups under Welcome and Next Steps, and the main session at 7.30, Antigravity for non-coders, where the hackathon groups would be revealed.

HOW TO USE THESE NOTES

Resources promised for the platform: the compiled toolkit and repository guide, the half-hour video on connecting Gmail to a local n8n through Google Cloud Console, and the video-generation workflow template. n8n's own template library, over twelve thousand ready workflows filterable by role and software, was linked in chat.

Two habits Samruddhi asked everyone to adopt: keep a Changes Made sheet whenever an agent touches finances, so every row and formula it altered is logged, and add a filter or guardrail node at the front of any workflow that can send or update, so only trusted senders pass through.

Designs and Agents on a Live Website

The opening question bundled two that need different answers. Putting a finished Lovable design on a live website is a domain question: when you publish, Lovable offers a custom domain, and the registrar, Hostinger, GoDaddy, or Namecheap, gives you DNS records under the domain's settings that you paste in the way you paste an API key, after which the design serves at your address. Putting an n8n agent on a live website is a webhook question: the workflow starts with a webhook node configured on the site, in vibe coding by telling Lovable itself to send a certain kind of data to that webhook when users interact, and ends with a Respond to Webhook node only if the site needs to show the output, as a chatbot must, but not if the data is merely being stored. The follow-up about a marketplace of reusable workflows got n8n's official template library, some twelve and a half thousand automations filterable by role and tool, to build on rather than from scratch. Adding voice to Jerry was traced to the trigger: pick a platform that accepts audio, WhatsApp or Telegram, branch on message type, download and transcribe the audio, and pass text to the agent, because models can analyse audio directly but at far higher cost. Smaller threads went quickly: PDFs become Markdown through any free online converter; a stubborn Cursor-to-Bitbucket connection is handed to Claude Desktop with desktop use, which researches the errors others reported and takes the steps; AutoCAD and Revit files are too large for n8n, Claude, or ChatGPT to process directly, so those tools convert geometry into code before their own AI reads it, a layer you would have to add; custom GPTs, on paid plans, take a name, attached files, and instructions, then publish privately, by link, or publicly; n8n's assistant can build whole workflows from plain language; and any JSON, published or not, exports from the three-dot menu.

Claude Code, Licences, and the Privacy Question Again

Asked for the best harness of connectors and plugins for vibe coding in Claude Desktop, her answer was that you do not need one: switch from chat to Claude Code, available in the terminal and the desktop app, and it writes the code, opens its own browser tab through browser use, tests and debugs there while you do other work, and returns the result. Licences, raised after Dileep's session, live on the product's website or repository, shown live on n8n's GitHub page where the enterprise licence links to its full text; MIT is the safe default for anything you build, and clients rarely ask, since you are integrating into systems they already license, so create no licence unless required. The privacy worry about connecting Gmail, Slack, and Asana to Claude or ChatGPT was answered with Jerry on screen: connected does not mean read. Nothing happens until a workflow is triggered, and then only the specific tool, get many messages for the last five emails, runs, touching only that data; the same holds in Claude and ChatGPT, where a connector's individual tool fires only for the prompt that needs it. Two related answers rounded it out: yes, Claude Desktop or ChatGPT can observe your Excel and Gmail work for days through desktop use and turn it into skills, but keep one skill per role, Excel, email, and so on, under a master skill that delegates, exactly like the chief-of-staff bot from Grokbot; and for accountants and controllers, enterprise plans remove the privacy concerns, Ollama keeps everything local, and the Changes Made log protects you from an agent's silent mistake in a spreadsheet.

Zapier Versus n8n, FAL Pricing, and Architecture

Zapier is good, she said, because its specific nodes do the job for many people without anything more; n8n is the most customisable platform on the market and the only one you can download and host on a client's own server, which matters for security-sensitive clients. If Zapier does the job, stay; switch when customisation or self-hosting demands it, because the purpose of learning AI is getting the work done the fastest, easiest way. Video-generation cost has no single number, the variance between models is a hundredfold, so she opened FAL, which lists every creative model with a playground where you set resolution and duration and read the per-second or per-minute price beneath the output, no API connection needed to test. The architecture question, ten independent workflows or one, got the rule map the process first: if one flowchart is too complicated, draw several, and put an orchestrator workflow above them that triggers each as needed, the pattern of Jerry with its three agents and of Jarvis with six workflows behind one master; where MCP servers are unnecessary, n8n's Execute Workflow node does the same job as a tool or a trigger. For newcomers who lack process mapping, Grokbot needs none, it asks its own questions, while n8n rewards a plain flowchart, marketing as research, generation, and analytics, each with sub-processes, each box then assigned a node or a tool.

Three Ways to Get Paid

Monetising n8n work comes in three models. A one-time build, where the client describes the process, you build it, hand over the JSON, and charge once. Build plus maintenance, her recommended default, because in ninety-nine percent of cases the first build reflects what the client said in meetings, and real use demands customisation, new models arrive daily and the workflow should stay optimised, and dependencies break, so the monthly fee covers agreed customisations, model updates, and uptime through error workflows and model switching. And subscription, for one workflow with a front end that many people use repeatedly. Whose n8n runs it depends on the client: some want it on their own servers, which you set up, and some accept your account, where the enterprise-style projects feature keeps each client in a separate project whose credentials are invisible to every other workspace. API keys she prefers to own, one per client and even per department where billing is separate, because a client's own key hides what is happening, how many tokens, which model, whether a cheaper one would do, while your key shows all of it and text-model costs are minimal enough to pass through later. How she found clients was the last of these: posting and waiting did not work, so she went to her network, who already trusted her in a field she was simply adding AI to; the inexperienced should build projects, post them, assemble a portfolio, bid on freelancing platforms, launch on Product Hunt, and tell their communities; and cold outreach works only after defining the workflow and the ideal client profile, never sprayed at everyone.

API, Webhook, MCP, and Hybrid RAG

The MCP-versus-API question got the trio in order. APIs came first: an API lets one piece of software talk to another, an API key is the secret that permits it, and an endpoint is one specific action, as Jarvis's calendar tools, check availability, get events, delete, update, reschedule, are each separate endpoints underneath; you must call an API, it never volunteers. Webhooks are always live and fire when something happens, as Retell's call-completion webhook did when it pushed the summary into Google Sheets. MCPs, launched only a year or two ago with agents, are a layer between the agent and the APIs, impossible without them, that reduces hallucination and shrinks context, because twenty tools wired directly to one agent become a few MCP servers that each route to the right call. Related answers followed: Meta's ad APIs and the campaign templates in n8n can build dashboards and alert you when KPIs cross thresholds you define; an animated on-screen character that takes input and runs a workflow is a vibe-coding prompt in Claude Code or Antigravity, not a special product; and an e-commerce support bot needs a RAG, specifically a hybrid RAG, keyword search to match an exact order number or email against the CRM and vector search to answer from policies, the two combined so a refund question finds both the customer and the rule. The coloured groupings in her workflows are sticky notes, added from the note icon, recoloured, resized, and labelled.

Avatars, Storyboards, Guardrails, and When to Leave n8n

A HeyGen avatar that felt unique but not production-ready got the practitioner's checklist: expressions and lighting matter most, the training voice should be natural speech rather than reading, gestures and body language are captured exactly as performed, so record several poses and styles, seated, standing, walking, as the platform's own stock avatars do, and keep multiple avatars for multiple situations; Google Flow now creates avatars too. Image and video models were named again, Higgsfield for both, HeyGen for avatars, FAL for the whole catalogue, ElevenLabs for Indian-language voices. Story videos with many characters begin with a storyboard, each frame numbered, handed to an image-to-video model that stitches the series, shown on her own chocolate-advertisement board in Flow. Guardrails against jailbreaks on agents with send or update permissions come in two layers, n8n's own guardrail nodes configured per risk, and source filtering, her Telegram workflow passing only messages from her own username. Moving from n8n cloud to local is an export of all workflows from the admin panel; Yahoo or any custom mail connects through the generic email node with server URLs; and if ChatGPT does not appear in Ollama's app list, the ChatGPT desktop app is probably not installed. The closing answer settled the week's recurring confusion, n8n or Codex and Claude: everything n8n does can be done in code, but n8n shows every node and its configuration in drop-downs a non-coder can read, whereas a coded workflow is opaque in a language you may not recognise and hard to debug when it breaks, so stay in n8n unless you can vibe code, and use Codex or Claude Code when you can.

THE MENTAL MODEL

A domain connects a design and a webhook connects an agent. Connected is not read; nothing runs until it is triggered, and then only the one tool the prompt needs. Map the process, orchestrate several workflows under one master, charge for build plus maintenance on your own API keys in per-client projects, and reach for MCPs to cut hallucination, hybrid RAG to find both the customer and the rule, and guardrails at the front door of anything that can send.

Key Takeaways

Stripped to its essentials, the hour made the following points.

- **Designs need a domain, agents need a webhook.** Paste your registrar's DNS records into Lovable's custom-domain option, start the n8n workflow with a webhook the site calls, and add Respond to Webhook only when the site must show the answer.
- **Start from twelve thousand templates.** n8n's library filters by role and tool, its assistant builds workflows from plain language, and voice arrives through a Telegram or WhatsApp trigger transcribed before the agent.
- **Claude Code needs no harness for vibe coding.** Switch from chat to code and it writes, tests in its own browser tab, and debugs while you work; licences live on the product's site, and MIT is the default for what you build.
- **Connected is not read.** Authorising Gmail or Slack does nothing until a prompt triggers one specific tool on one specific request, in n8n, Claude, and ChatGPT alike.
- **Stay in Zapier until you need n8n.** Zapier's nodes suffice for many; n8n wins on customisation and on self-hosting for security-sensitive clients, and the only rule is getting the job done fastest.
- **Charge for build plus maintenance.** First builds always need customising, models change daily, and dependencies break; keep clients in separate n8n projects with your own per-client API keys so usage and cost stay visible.
- **Know the trio and the hybrid.** APIs are called, webhooks fire, MCPs route agents to APIs with less hallucination and context; e-commerce support needs hybrid RAG, keyword search for the order and vector search for the policy.
- **Guard the front door and log the changes.** Filter senders and configure guardrail nodes on anything that can send or update, keep a Changes Made sheet for finance work, and start client-hunting with the network that already trusts you.

End of Office Hours Six notes. Next: Office Hours Seven at 6 pm IST.