

AI GENERALIST ACCELERATOR · 14-DAY AI SPRINT

Bonus Office Hours Seven

Office Hours Seven: The Toolkit, the Build Ladder, and an AI Roadmap for a Company

A cleaned, detailed record of the seventh Slido question hour

Topics covered: Attachments, Voice, and Contacts for Jerry · The AI Toolkit Website · Scrapers, Numbers, and Memory Layers · Antigravity Credits and Global Skills · The Three-Tier Build Ladder · Multi-Tenant Tools for Clients · Licences, Security, and Privacy · An AI Roadmap for a Company · First Clients and Pricing

About This Session

The seventh office hour opened day eleven at 6 pm IST with Sumedha hosting and Samruddhi Mokal answering for the fourth day running, at a brisker pace than usual because the hackathon was starting that night and the Slido queue was long. The hour's centrepiece was a delivery: the AI toolkit website Samruddhi had promised all week, categorised from chat assistants to grants, was shared in chat and toured live, with the promise to keep adding to it. A learner's request, that every obscure tool she recommends come with a link because the captions mangle the names, was granted on the spot. The recurring boundary held: programme, workbook-support, bonus-tool, and recording questions were routed to Sumedha, who at the close pointed everyone to the Ask Doubt channel, where a mentor named Santesh replies within twenty-four hours, and to the group allocator, where fresh group allocations had just gone live for anyone who had raised a request. The main session, Claude mastery with Vishwas followed by the hackathon step-two briefing, began after a thirty-minute break.

HOW TO USE THESE NOTES

The AI toolkit website is the resource of the week: categories for chat assistants, model access and routing, scraping, image and audio, voice agents and phone numbers, app builders, coding assistants, sales and outreach, social media, meeting notes, presentations, whiteboarding, backends, payments, hosting, GPUs, MCPs and skills, security, memory layers, communities, free credits, accelerators, grants, and autonomous-agent repositories. Its link is in the chat and heading for the office-hours resources.

Two habits Samruddhi asked for: check your Antigravity model quotas under Settings, Models, before a long build, and never put credit-card or financial details into any tool while testing; email access is fine because permissions do nothing until a workflow is triggered, and can be revoked before that.

Attachments, Voice, and Contacts for Jerry

The first three questions extended Jerry. An agent that reads emails with PDF attachments and reconciles them into an Excel sheet uses the Gmail node's own options, every n8n node hides filters and extras behind its more-options panel, including attachment download, after which extracted data flows to the sheet node and invoices are created or checked through the invoicing tool's API, one endpoint to create, another to check payment status. Voice commands replace the chat trigger with a platform that accepts audio, Telegram or WhatsApp, and n8n allows several triggers on one workflow, so text and voice can both feed it, with a classification step that downloads and transcribes audio before the agent, because passing audio straight to a model burns credits. Jerry asking for email addresses every time is solved by a Google Contacts node, as her own Jarvis has, so a name resolves to an address, with filters for anything else the send needs. Transcription models are found by searching the node list, hers being OpenAI's; Yahoo or any provider without a node uses the generic email node, or n8n's assistant builds the integration when asked; drop-downs in a Google Sheet are set in the sheet-creation node's filters; sticky notes are copied by double-clicking and using normal keyboard shortcuts; and one workflow calls another through the when-executed-by-another-workflow trigger, with the reminder never to congest everything into one canvas and always to attach an error workflow.

The Toolkit Tour

The website walk was fast and opinionated. Under model access, alongside OpenRouter, FAL, and Higgsfield, she added Groq, spelled with a Q and distinct from xAI's model, a router-like service that gives free credits every day and has its own n8n chat-model node. Under scraping, Apify remains the marketplace of thousands of scrapers, chosen by user count and rating, with one rule for social platforms: pick scrapers that do not use cookies, because cookies tell the site it is you, and a mock account in incognito is a fair alternative, illustrated by opening a page's inspector to show what a scraper actually reads. Voice agents got an India-specific answer to a week of questions, Bolna and other Indian providers for building, and Wobbies for phone numbers starting with plus ninety-one, integrated with Vapi, ElevenLabs, Retell, and the rest. Among app builders, Base44 was singled out for mobile apps because it builds to App Store and Play Store guidelines and integrates with both. For cold outreach, Instantly warms inboxes, sends, tracks opens and clicks, and runs follow-up sequences, with Smartlead similar; for social media, PhantomBuster, an n8n-like workflow tool built for LinkedIn, Instagram, Twitter, and HubSpot, exists because raw n8n bots get accounts blocked, with pre-built flows for extracting leads from posts, connecting, and following up. Miro is her whiteboard; Supabase, Firebase, and Pinecone cover backends and vectors; Stripe and Razorpay payments; Vercel and Netlify hosting; RunPod rents GPUs by the hour for models a laptop cannot carry. Memory layers earned a personal verdict, Mem0 and SuperMemory demand configuration while Memory, her daily driver, reaches a thousand-plus apps, multiple Gmail and calendar accounts, and takes voice notes in several languages over WhatsApp. The tail of the site holds communities, JoinSecret for limited free credits and coupons, accelerators and fellowships for funding, grants for money nobody expects back, and a list of autonomous-agent GitHub repositories.

Antigravity Credits, Skills, and the Build Ladder

Running out of Antigravity credits got a prevention habit rather than a workaround: Settings, then the user settings, then Models, where every model shows weekly and hourly remaining quota, so switch models before a wall rather than discovering it. Multiple projects mean one folder per project, and skills that should work across projects are installed at global rather than folder level simply by saying so in the prompt. Antigravity does not connect to n8n directly, a vibe-coding tool and an automation tool are different things; a front end built in Antigravity reaches n8n through n8n's MCP server URL. Then the ladder she asked everyone to memorise. Tier one, a beautiful front end: Lovable. Tier two, a decent backend: Bolt, Replit, or Google's AI Studio. Tier three, a heavy backend and a solid structure on both sides: the IDEs, Antigravity, Cursor, and her own choice, Claude Code. Tiers convert: build the MVP in Lovable, connect it to GitHub when real users arrive, and tell Antigravity to pull the repository and build on top, which is also the answer to the Antigravity-versus-ChatGPT question, ChatGPT is an agent that can add connectors, Antigravity and Codex are coding tools that create, connect, and deploy front ends and back ends, with Codex being OpenAI's equivalent. Lovable versus Wix: Lovable builds web apps, mobile apps, and websites, Wix only websites. Cloning a site you admire is a native strength of every vibe-coding tool given the link, as her Spotify-like Replit build showed; LangChain is n8n-like automation for the technically inclined; and Google Flow edits existing footage for lighting and background.

Tools for Clients, Governance, and a Company Roadmap

Serving many small-business clients with one product is a structure question: one website, multi-user authentication, a Supabase backend recording each user's actions, and three permission layers, you as super admin, the client's CEO as admin who sets their team's permissions, and the team as sub-accounts, all added by prompt while vibe coding. Building client automations faster in n8n has two routes, n8n's MCP connected to Claude, or n8n's own assistant, which produces a random workflow from a vague prompt and a good one from a specific brief that names the tools and actions. Reminder systems for payments and tasks run a scheduled trigger against the CRM each morning, filter by date, pull the contact, and branch to Retell for calls, Twilio for SMS, and a WhatsApp provider for messages. Governance follows the client's policies, not yours: workflows on their servers if required, and Gmail without OAuth means a Google Cloud Console web app, client secret, and scoped permissions, an AI layer added on top of rules that already exist. The hundred-person, four-location roadmap got her method: name the single goal, revenue, cost, or scale without headcount; find the department that reaches it fastest, sales for revenue; automate one process end to end, ICP to scraping to research to a hyper-personalised message to send; track the KPI, iterate or move on; and above all automate processes that already exist on tools already in use rather than inventing new ones, starting from a plain-English flowchart. Accounting processes, reconciliations, and budget variances follow the same rule, plugins or connectors for the software in use, or its API keys inside an n8n workflow that mirrors the step-by-step process.

Licences, Privacy, and Learning Without Drowning

Apache joined the licence lesson with a robot analogy: MIT is your friend saying use my robot however you like; Apache is the same plus explicit access to the secret that makes it walk, which MIT permits in practice but does not spell out, leaving a theoretical legal opening; GPL is set aside for anything commercial, since it demands everything be open-sourced. Privacy questions came in three forms. Testing new tools: never share financial or card details, personal email is fine, read the permissions, and remember that authorisation does nothing until a workflow runs. Does n8n store your emails when it reads them: it does not store them, but the content is passed to whichever LLM node processes it, exactly as prompting ChatGPT with a connector does. Does RunPod defeat local privacy: the data leaves your machine for a dedicated server for that session, as enterprises do on AWS, but it is not passed to the model provider, you are renting hardware, not sharing data. HeyGen avatars do not self-select; each pose and tone is a separately recorded avatar with its own voice and video IDs, chosen per situation in the workflow. And to the learner asking how she got so good, and how to stop feeling overwhelmed: years of experience from before GPT could do more than complete a sentence, the right communities and people, and a discipline she wished she had started with, pick one use case, automate or vibe-code it until it shows immediate impact on your life, then move to the next, because time spent surveying tools instead of shipping something is where confidence dies. The why-do-all-these-platforms-exist question got the same answer, they are competitors, and you stick to whichever gets your job done fastest.

First Clients, Pricing, and Hosting

Her first client came through her existing network, told in plain terms: she announced what she was doing, asked contacts what problems they faced, offered to solve one with AI, built a portfolio from the results, and shared her work on social media until traction came, and today most clients arrive by referral, so she posts less. Cold outreach and social media both work, but build your own pipeline. India or the West is a personal choice; start wherever the first client comes fastest, collect testimonials, then diversify. Pricing has no range, only inputs: your ideal client profile, their problem statement, the return on investment you deliver, and what they already pay to solve it, enterprises high and small businesses to their means. Hosting for herself is Vercel, free until the backend gets heavy, with Railway as the cheap step beyond it and client work living on the client's own servers. A one-card-for-all-tools service does not exist and should not be trusted, and SEO automation exists as ready n8n templates, her trend and cluster workflows finding high-volume keywords, choosing topics, writing, internally linking, adding cover images, and publishing to WordPress or anywhere, with the SEO vocabulary itself best learned by asking Claude.

THE MENTAL MODEL

Every node has hidden options; every provider without a node has an API or the generic node. Route models through Groq or OpenRouter, scrape without cookies, remember with a memory layer, and climb the build ladder from Lovable to Bolt to the IDEs by way of GitHub. Serve clients with one multi-tenant product under their governance, roadmap a company from a single goal through the fastest department on tools it already uses, and grow your own confidence one shipped use case at a time.

Key Takeaways

Stripped to its essentials, the hour made the following points.

- **Jerry grows through node options.** Attachments come from the Gmail node's extras, voice from a Telegram or WhatsApp trigger transcribed before the agent, addresses from a Google Contacts node, and any missing provider from the generic email node or the assistant.
- **The toolkit is live and categorised.** Groq for free daily model credits, Apify without cookies, Bolna and Wobbies for Indian voice agents and numbers, Base44 for mobile, Instantly and PhantomBuster for outreach, RunPod for GPUs, Memory for a memory layer, JoinSecret for credits, and grants for money nobody expects back.
- **Check quotas and set skills globally.** Antigraivty's Settings, Models panel shows hourly and weekly limits per model, one folder per project, and a global-level instruction in the prompt shares skills across projects.
- **Climb the ladder through GitHub.** Lovable for front ends, Bolt, Replit, or AI Studio for a decent backend, Antigraivty, Cursor, or Claude Code for heavy ones, and an MVP moves up by connecting Lovable to GitHub and pulling the repository into the IDE.
- **One product, three permission layers.** Multi-user auth on a Supabase backend with super admin, client admin, and sub-accounts serves many clients from one link, under each client's own governance and servers.
- **Know MIT from Apache from GPL.** MIT lets you use the robot, Apache hands over the secret that makes it walk, GPL forces everything open, so keep GPL away from anything you sell.
- **Roadmap from one goal.** Name revenue, cost, or scale, automate the fastest department's existing process on existing tools from a plain-English flowchart, track the KPI, then move to the next department.
- **Ship one use case, then the next.** Confidence and competence come from immediate impact, not tool surveys; first clients come from your network and referrals, and prices come from the client's ICP, problem, and current spend.

End of Office Hours Seven notes. Next: Office Hours Eight at 6 pm IST.